

EASY Meta-programming with Rascal

Paul Klint, Tijs van der Storm, and Jurgen Vinju
Centrum Wiskunde & Informatica
and
Universiteit van Amsterdam

Abstract. Rascal is a new language for meta-programming and is intended to solve problems in the domain of source code analysis and transformation. In this article we give a high-level overview of the language and illustrate its use by many examples. Rascal is a work in progress both regarding implementation and documentation. More information is available at <http://www.rascal-mpl.org/>.

Acknowledgement. This project was carried out as part of ATEAMS, a joint research project team funded by CWI (Netherlands) and INRIA (France).

Keywords: source code analysis, source code transformation, meta-programming, domain-specific languages.

1 A New Language for Meta-programming

Meta-programs are programs that analyze, transform or generate other programs. Ordinary programs work on data; meta-programs work on programs. The range of programs to which meta-programming can be applied is large: from programs in standard languages like C and Java to domain-specific languages for describing high-level system models or applications in specialized areas like gaming or finance. In some cases, even test results or performance data are used as input for meta-programs. Rascal is a new language for *meta-programming*, this is the activity of writing meta-programs.

1.1 The EASY Paradigm

Many meta-programming problems follow a fixed pattern. Starting with some input system (a black box that we usually call *system-of-interest*), first relevant information is extracted from it and stored in an internal representation. This internal representation is then analyzed and used to synthesize results. If the synthesis indicates this, these steps can be repeated over and over again. These steps are shown in Figure 1.1, “EASY: the Extract-Analyze-Synthesize Paradigm”.

This is an abstract view on solving meta-programming problems, but it has wide applicability. Let's illustrate it with a few examples.




Fig. 1.1. EASY: The Extract-Analyze-Synthesize Paradigm

1.1.1 Finding Security Breaches

Alice is system administrator of a large online marketplace and she is looking for security breaches in her system. The objects-of-interest are the system's log files. First relevant entries are extracted. This will include, for instance, messages from the SecureShell demon that reports failed login attempts. From each entry, login name and originating IP address are extracted and put in a table (the internal representation in this example). These data are analyzed by detecting duplicates and counting frequencies. Finally results are synthesized by listing the most frequently used login names and IP addresses.

1.1.2 A Forensic DSL Compiler

Bernd is a senior software engineer working at a forensic investigation lab of the German government. His daily work is to find common patterns in files stored on digital media that have been confiscated during criminal investigations. Text, audio and video files are stored in many different data formats and each data format requires its own analysis technique. For each new investigation, ad hoc combinations of tools are used. This makes the process very labour-intensive and error-prone. Bernd convinces his manager that designing a new domain-specific language (DSL) for forensic investigations may relieve the pressure on their lab. The DSL should at least be able to define multiple data formats and generate recognizers for them. After designing the DSL---let's call it DERRICK---he makes an EASY implementation for it. Given a DERRICK program for a specific case under investigation, the DERRICK compiler first extracts information from the program and analyzes it further: which media formats

are relevant? Which patterns to look for? How should search results be combined? Given this information, Java code is synthesized that corresponds to the meaning of the DERRICK program: specialized generated code that glues together various existing libraries and tools.

1.1.3 Renovating Financial Software

Charlotte is software engineer at a large financial institution and she is looking for options to connect an old and dusty software system to a web interface. She will need to analyze the sources of that system to understand how it can be changed to meet the new requirements. The objects-of-interest are in this case the source files, documentation, test scripts and any other available information. They have to be parsed in some way in order to extract relevant information, say the calls between various parts of the system. The call information can be represented as a binary relation between caller and callee (the internal representation in this example). This relation with 1-step calls is analyzed and further extended with 2-step calls, 3-step calls and so on. In this way call chains of arbitrary length become available. With this new information, we can synthesize results by determining the entry points of the software system, i.e., the points where calls from the outside world enter the system. Having completed this first cycle, Charlotte may be interested in which procedures can be called from the entry points and so on and so forth. Results will be typically represented as pictures that display the relationships that were found. In the case of source code analysis, a variation of our workflow scheme is quite common. It is then called the extract-analyze-view paradigm and is shown in Figure 1.2, “The extract-analyze-view paradigm”.

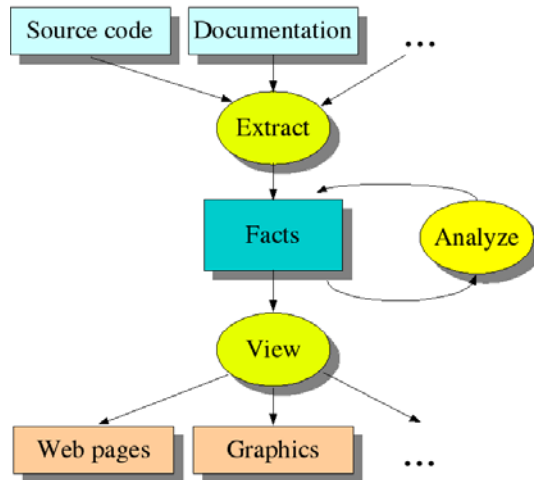


Fig. 1.2. The extract-analyze-view paradigm

1.1.4 Finding Concurrency Errors

Daniel is concurrency researcher at one of the largest hardware manufacturers worldwide. Concurrency is the big issue for his company: it is becoming harder and

harder to make CPUs faster, therefore more and more of them are bundled on a single chip. Programming these multi-core chips is difficult and many programs that worked fine on a single CPU contain hard to detect concurrency errors due to subtle differences in the order of execution that results from executing the code on more than one CPU. He is working on tools for finding concurrency errors. First he extracts facts from the code that are relevant for concurrency problems and have to do with calls, threads, shared variables and locks. Next, he analyzes these facts and synthesizes an abstract model that captures the essentials of the concurrency behaviour of the program. Finally he runs a third-party verification tool with this model as input to do the actual verification.

1.1.5 Model-Driven Engineering

Elisabeth is a software architect at a large airplane manufacturer and her concern is reliability and dependability of airplane control software. She and her team have designed a UML model of the control software and have extended it with annotations that describe the reliability of individual components. She will use this annotated model in two ways: (a) to extract relevant information from it to synthesize input for a statistical tool that will compute overall system reliability from the reliability of individual components; (b) to generate executable code that takes the reliability issues into account.

1.2 Rascal

With these examples in mind, you have a pretty good picture how EASY applies in different use cases. All these cases involve a form of *meta-programming*: software programs (in a wide sense) are the objects-of-interest that are being analyzed, transformed or generated. The Rascal language you are about to learn is designed for meta-programming following the EASY paradigm. It can be applied in domains ranging from compiler construction and implementing domain-specific languages to constraint solving and software renovation.

Since representation of information is central to the approach, Rascal provides a rich set of built-in data types. To support extraction and analysis, parsing and advanced pattern matching are provided. High-level control structures make analysis and synthesis of complex data structures simple.

1.3 Benefits of Rascal

Before you spend your time on studying the Rascal language it may help to first hear our elevator pitch about the main benefits offered by the language:

- **Familiar syntax** in a *what-you-see is-what-you-get* style is used even for sophisticated concepts and this makes the language easy to learn and easy to use.
- **Sophisticated built-in data types** provide standard solutions for many meta-programming problems.

- **Safety** is achieved by finding most errors before the program is executed and by making common errors like missing initializations or invalid pointers impossible. *At the time of writing, this checking is done during execution but a static typechecker is nearing completion.*
- **Local type inference** makes local variable declarations redundant.
- **Pattern matching** can be used to analyze all complex data structures.
- **Syntax definitions** make it possible to define new and existing languages and to write tools for them.
- **Visiting** makes it easy to traverse data structures and to extract information from them or to synthesize results.
- **Templates** enable easy code generation.
- **Functions as values** permit programming styles with high re-use.
- **Generic types** allow writing functions that are applicable for many different types.
- **Eclipse integration** makes Rascal available in a familiar environment: all familiar tools are at your fingertips.

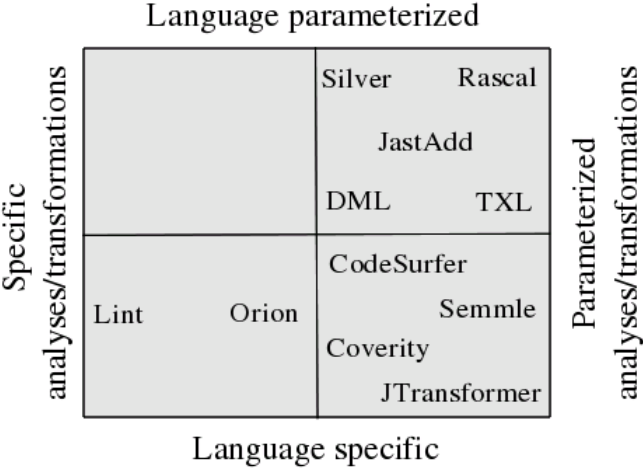


Fig. 1.3. Positioning of various systems

1.4 Related Work

We discuss related tools and specific influences of earlier work on the Rascal language. Figure 1.3, “Positioning of various systems” shows the relative positioning of various systems to be discussed in this section. On the vertical axis, the level of language-

specificity is shown. Systems at the bottom of the figure are language-specific and can only be applied to programs in one specific language. Systems at the top of the figure are language-parameterized, i.e., they are generic and can be applied to arbitrary languages. On the horizontal axis, the level of parameterization of analysis and transformation is displayed. Systems on the left have fixed analysis/transformation properties while the systems more to the right have more and more support for defining analysis and transformation operations.

1.4.1 Related Tool Categories

We focus on positioning Rascal relative to a number of broad categories of tools that are available and give some examples of typical tools in those categories.

Lexical Tools. Lexical tools are based on the use of regular expressions and are supported by widely available systems like Grep, AWK, Perl, Python and Ruby. Regular expressions have limited expressiveness since they cannot express nested structures. Lexical tools are therefore mostly used to extract unstructured textual information, like counting lines of code, collecting statistics about the use of identifiers and keywords, and spotting simple code patterns. Experience shows that maintaining larger collections of regular expressions is hard.

Compiler Tools. Well-known tools like Yacc, Bison, CUP, and ANTLR [Par07] fall in this category. They all use a context-free grammar and generate a parser for the language defined by that grammar. Some provide automation for syntax tree construction but all other operations on parsed programs have to be programmed manually in a standard programming language like C or Java.

Attribute Grammar Tools. An attribute grammar is a context-free grammar enriched with attributes and attribute equations that express dependencies between attributes. In this way semantic operations like typechecking and code generation can be expressed at a high level. Tools in this category are FNC-2 [JPJ+90], JastAdd [HM03], Silver [VWBGK10], and others. Attribute Grammar tools are strong in program analysis but weaker in program transformation.

Relational Analysis Tools. Relational algebra tools provide high-level operations on relational data. They are strong in expressing analysis problems like, for instance, reachability. Examples of tools in this category are GROK [Hol08], Crocopat [Bey06] and RScript [Kli08]. These tools assume that the data are already in relational form and provide support for neither extraction nor transformation.

Transformation Tools. Tools in this category focus on transforming programs from one representation to another representation and they are frequently based on term rewriting. Examples are ASF+SDF [Kli93], [BDH+01], ELAN [BKK+98], STRATEGO [BKVV08], TOM [BBK+07], and TXL [Cor06].

Logic-based Tools. Logic-based languages/tools form a category on their own since they partially overlap with the other categories. The primary example is Prolog that

has been used for syntax analysis, analysis and synthesis and is still in use today. [KHR07] gives an overview of various logic-based approaches. Nonetheless, these tools have never come into wide use. [LR01] illustrates the use of Prolog for analysis and transformation. [KHR07] describes JTransformer, a more recent Prolog-based transformation tool for Java.

Miscellaneous Tools. There is a wide range of tools available in the space sketched in Figure 1.3, “Positioning of various systems”. Lint [Joh79] is the classical example of a command-line tool that performs C-specific analysis. ORION [DN04] is another example of a tool for fixed error analysis for C/C++ programs. Codesurfer [AZ05] and Coverity [BBC+10] are examples of tools for the analysis of C programs that provide programmable analyses. Semmle [dMSV+08] is an example of a Java-specific tools with highly-customizable queries. The same holds for JTransformer [KHR07], but this also supports customizable transformations.

Comparing Tool Categories. We audaciously list the strengths and weaknesses of these tool categories in Table 1.1, “Comparing Tool Categories”. There are large variations among specific tools in each category but we believe that this table gives at least an impression of the relative strengths of the various categories and also documents the ambition we have with the Rascal language. It is clear from this overview that the level of integration of extraction, analysis and synthesis is mostly weak. Exceptions to this are systems that explicitly integrate these phases. For language-specific systems this are, for example, CodeSurfer [AZ05], Coverity [BBC+10], Semmle and JTransformer. For language-parametric systems examples are DMS, TXL, JastAdd and Silver. All these systems differ in the level of integration and IDE support. Rascal is comparable to the latter systems regarding language-parametricity, provides good IDE support and outperforms many other systems with respect to the integrated and advanced mechanisms for specifying analysis and transformation.

Table 1.1. Comparing Tool Categories

Type	Extract	Analyze	Synthesize
Lexical Tools	++	+ / -	--
Compiler Tools	++	+ / -	+ / -
Attribute Grammar Tools	++	+ / -	--
Relational Tools	--	++	--
Transformation Tools	--	+ / -	++
Logic-based Tools	+ / -	+	+
<i>Our Goal: Rascal</i>	++	++	++

1.4.2 Work Directly Related to Rascal

Rascal owes a lot to other meta-programming languages, in particular the user-defined, modular, syntax and term rewriting of ASF+SDF [Kli93], [BDH+01], the relational

calculus as used in RScript [Kli08] and pioneered by GROK [Hol08], traversal functions as introduced in [BKV03], strategies as introduced in ELAN [BKK+98] and Stratego [BKVV08], and integration of term rewriting in Java as done in TOM [BBK+07].

We also acknowledge less specific influences by systems like TXL [Cor06], ANTLR [Par07], JastAdd [HM03], Semmle [dMSV+08], DMS [BPM04], and various others like, for instance, ML and Ruby.

The application of Rascal for refactoring is described in [KvdSV09]. Using Rascal for collecting source code metrics for DSL implementations is presented in [KvdSV10].

1.4.3 Contribution

We claim as main contribution of Rascal the seamless integration of extraction, analysis and synthesis concepts into a relatively simple, statically typed, language with familiar syntax. Specific contributions are the integration of syntax definitions and parsing functions, access to syntax trees both as fully typed and as untyped data structures, the rich collection of data types and patterns, string templates, rewrite rules as normalization device for structured data, and the use of transactions to undo side-effects. All these features are discussed in some detail in this paper.

1.5 Reading Guide

The aim of this article is to give an easy to understand but comprehensive overview of the Rascal language and to offer problem solving strategies to handle realistic problems that require meta-programming. Problems may range from security analysis and model extraction to software renovation, domain-specific language development and code generation.

The scope of this article is limited to the Rascal language and its applications but does not address implementation aspects of the language.

The structure of the description of Rascal is shown in Figure 1.4, “Structure of the Rascal Description”. This article provides a self-contained version of the first three parts:

- **Introduction:** gives a high-level overview of Rascal and consists of Section 1, “A New Language for Meta-Programming” and Section 2, “Rascal Concepts”. It also presents some simple examples in Section 3, “Some Simple Examples”.
- **Problem Solving:** describes the major problem solving strategies in Rascal’s application domain, see Section 4, “Problem Solving Strategies”.
- **Examples:** gives a collection of larger examples, see Section 5, “Larger Examples”.

The other parts can be found online at <http://www.rascal-mpl.org>:

- **Reference:** gives a detailed description of the Rascal language, and all built-in operators and library functions.

- **Support:** gives tables with operators and library functions, a bibliography and a glossary that explains many concepts that are used in the descriptions of Rascal and tries to make them self-contained.

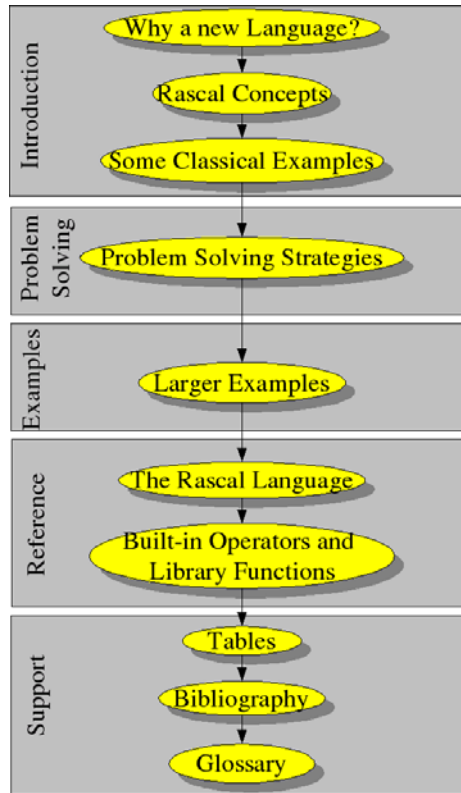


Fig. 1.4. Structure of the Rascal Description

1.6 Typographic Conventions

Rascal code fragments are always shown as a listing like this:

```
.. here is some Rascal code ...
```

Interactive sessions are shown as a screen like this:

```
rascal> Command;  
Type: Value
```

where:

- `rascal>` is the prompt of the Rascal system.
- **Command** is an arbitrary Rascal statement or declaration typed in by the user.

- **Type:** Value is the type of the answer followed by the value of the answer as computed by Rascal. In some cases, the response will simply be `ok` when there is no other meaningful answer to give.

2 Rascal Concepts

Before explaining the Rascal language in more detail, we detail our elevator pitch a bit and give you a general understanding of the concepts on which the language is based.

2.1 Values

Values are the basic building blocks of a language and the type of a value determines how they may be used.

Rascal is a value-oriented language. This means that values are immutable and are always freshly constructed from existing parts; sharing and aliasing problems are completely avoided. The language also provides variables. A value can be associated with a variable as the result of an explicit assignment statement: during the lifetime of a variable different (immutable) values may be assigned to it. Other ways to associate a value with a variable are by way of function calls (binding of formal parameters to actual values) and as the result of a successful pattern match.

2.2 Data Types

Rascal provides a rich set of data types:

- Booleans (`bool`).
- Arbitrary precision integers (`int`), reals (`real`) and numbers (`num`).
- Strings (`str`) that can act as templates with embedded expressions and statements.
- Source code locations (`loc`) based on an extension of Universal Resource Identifiers (URI) that allow precise description of text areas in local and remote files.
- Lists (`list`).
- Tuples (`tuple`).
- Sets (`set`).
- Maps (`map`).
- Relations (`rel`).
- Untyped tree structures (`node`)
- Fully typed data structures (`data`).

The elements of tuples, maps and relation may optionally be labelled for later reference and retrieval. There is a wealth of built-in operators and library functions available on the standard data types.

Table 1.2. Basic Rascal Types

Type	Examples
bool	true, false
int	1, 0, -1, 123456789
real	1.0, 1.0232e20, -25.5
str	"abc", "first\nnext", "result: <X>"
loc	file:///etc/passwd
datetime	\$2010-07-15T09:15:23.123+03:00
tuple[T_1, \dots, T_n]	<1,2>, <"john", 43, true>
list[T]	[], [1], [1,2,3], [true, 2, "abc"]
set[T]	{}, {1,2,3,5,7}, {"john", 4.0}
rel[T_1, \dots, T_n]	{<1,2>, <2,3>, <1,3>}, {<1,10,100>, <2,20,200>}
map[T, U]	(), (1:true, 2:false), ("a":1, "b":2)
node	f(), add(x,y), g("abc", [2,3,4])

The basic Rascal data types are illustrated in Table 1.2, “Basic Rascal Types”. Some types have another type as argument, for instance, `list[int]` denotes a list of integers. In the table, T , T_i and U denote such type arguments.

These built-in data types are closely related to each other:

- In a list all elements have the same static type and the order of elements matters. A list may contain the same value more than once.
- In a set all elements have the same static type and the order of elements does not matter. A set contains an element only once. In other words, duplicate elements are eliminated and no matter how many times an element is added to a set, it will occur in it only once.
- In a tuple all elements (may) have a different static type. Each element of a tuple may have a label that can be used to select that element of the tuple.
- A relation is a set of tuples that all have the same static tuple type.
- A map is an associative table of (key, value) pairs. Key and value (may) have different static types and a key can only be associated with a value once

Untyped trees can be constructed with the built-in type **node**. User-defined algebraic data types (ADTs) allow the introduction of problem-specific types and are a subtype of **node**. A fragment of the abstract syntax for statements (assignment, if, while) in a programming language would look as follows:

```
data STAT = asgStat(Id name, EXP exp)
          | ifStat(EXP exp, list[STAT] thenpart,
                  list[STAT] elsepart)
          | whileStat(EXP exp, list[STAT] body)
          ;
```

Syntax trees that are the result of parsing source files are represented as a data type (Tree). Since all datatypes are a subtype of **node**, this allows the handling of parse trees both as fully typed ADTs and as untyped nodes (thus enabling generic operations on arbitrary parse trees).

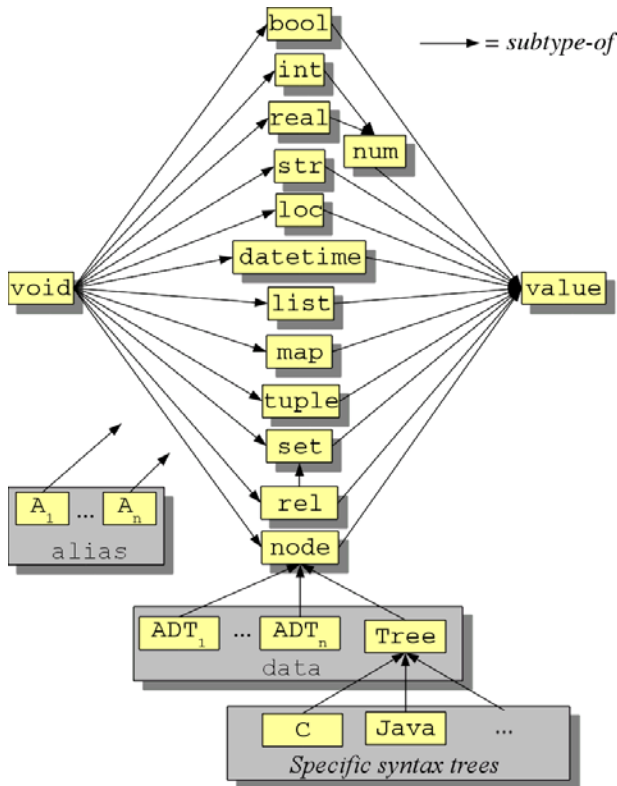


Fig. 1.5. Lattice of Rascal types

Without going into more details, we show the lattice of types in Figure 1.5, “Lattice of Rascal types”: all types are included between the minimal type (**void**) and the maximal type (**value**). The relation between **node**, abstract data type (**data**), **Tree** and parse trees

for different language is clearly shown. The alias mechanism allows shorthands for already defined types and is not further discussed.

2.3 Pattern Matching

Pattern matching determines whether a given pattern matches a given value. The outcome can be false (no match) or true (a match). A pattern match that succeeds may bind values to variables.

Pattern matching is *the* mechanism for case distinction (switch statement) and search (visit statement) in Rascal. Patterns can also be used in an explicit match operator `:` and can then be part of larger boolean expressions. Since a pattern match may have more than one solution, local backtracking over the alternatives of a match is provided. Patterns can also be used in enumerators and control structures like `for` and `while` statement.

A very rich pattern language is provided that includes string matching based on regular expressions, matching of abstract patterns, and matching of concrete syntax patterns. Some of the features that are provided are list (associative) matching, set (associative, commutative, idempotent) matching, and deep matching of descendant patterns. All these forms of matching can be used in a single pattern and can be nested. Patterns may contain variables that are bound when the match is successful. Anonymous (don't care) positions are indicated by the underscore (`_`).

Here is a *regular expression* that matches a line of text, finds the first alphanumeric word in it, and extracts the word itself as well as the text before and after it (`\W` matches all non-word characters; `\w` matches all word characters):

```
/^<before:\W*><word:\w+><after:.*$>/
```

Regular expressions follow the Java regular expression syntax with one exception: instead of using numbered groups to refer to parts of the subject string that have been matched by a part of the regular expression we use the notation:

```
<Name:RegularExpression>
```

If *RegularExpression* matches, the matched substring is assigned to string variable *Name*.

The following *abstract pattern* matches the abstract syntax of a while statement defined earlier:

```
whileStat(EXP Exp, list[STAT] Stats)
```

Variables in a pattern are either explicitly declared in the pattern itself---as done in the example, e.g., `EXP Exp` declares the variable `Exp` and `list[STAT] Stats` declares the variable `Stats`---or they may be declared in the context in which the pattern occurs. So-called *multi-variables* in list and set patterns are declared by a `*` suffix: `X*` is thus

an abbreviation for `list[...] X` or `set[...] X`, where the precise element type depends on the context. The above pattern can then be written as

```
whileStat(EXP Exp, Stats*)
```

where `Stats*` is an abbreviation for the declaration `list[STAT] Stats`. The choice whether `Stats` is a list or a set depends on the context. If you are not interested in the actual value of the statements the pattern can be written as

```
whileStat(EXP Exp, _*)
```

When there is a grammar for this example language (in the form of an imported SDF definition), we can also write *concrete patterns* as we will see below.

2.4 Enumerators

Enumerators enumerate the values in a given (finite) domain, be it the elements in a list, the substrings of a string, or all the nodes in a tree. Each value that is enumerated is first matched against a pattern before it can possibly contribute to the result of the enumerator. Examples are:

```
int x <- { 1, 3, 5, 7, 11 }
int x <- [ 1 .. 10 ]
/asgStat(Id name, _) <- P
```

The first two produce the integer elements of a set of integers, respectively, a range of integers. Observe that the left-hand side of an enumerator is a pattern, of which `int x` is a specific instance. The use of more general patterns is illustrated by the third enumerator that does a deep traversal (as denoted by the descendant operator `/`) of the complete program `P` (that is assumed to have a `PROGRAM` as value) and only yields statements that match the assignment pattern (`asgStat`). We have defined the `asgStat` constructor earlier in the example in Section 2.2, “Data Types” and here we see how it can be used as pattern. The descendant operator is part of the pattern and has as effect that the pattern is not only tried at the root of the subject to which it is applied (in this case `P`) but also to all its descendants, e.g., subtrees, list elements, and the like. Note the use of an anonymous variable at the `EXP` position in the pattern.

2.5 Comprehensions

Comprehensions are a notation inspired by mathematical set-builder notation that helps to write succinct definitions of lists and sets. They are also inspired by queries as found in a language like SQL.

Rascal generalizes comprehensions in various ways. Comprehensions exist for lists, sets and maps. A comprehension consists of an expression that determines the successive elements to be included in the result and a list of enumerators and tests (boolean expressions). The enumerators produce values and the tests filter them. A standard example is

```
{ x * x | int x <- [1 .. 10], x % 3 == 0 }
```

which returns the set {9, 36, 81}, i.e., the squares of the integers in the range [1 .. 10] that are divisible by 3. A more intriguing example is

```
{name | /asgStat(Id name, _) <- P}
```

which traverses program *P* and constructs a set of all identifiers that occur on the left hand side of assignment statements in *P*.

2.6 Control Structures

Control structures like *if* and *while* statements are driven by Boolean expressions, for instance

```
if(N <= 0)
  return 1;
else
  return N * fac(N - 1);
```

Actually, combinations of generators and Boolean expressions can be used to drive the control structures. For instance,

```
for(/asgStat(Id name, _) <- P, size(name) > 10){
  println(name);
}
```

prints all identifiers in assignment statements (*asgStat*) that consist of more than 10 characters.

2.7 Case Distinction

The switch statement as known from C and Java is generalized: the subject value to switch on may be an arbitrary value and the cases are arbitrary patterns followed by a statement. Here is an example where we take a program *P* and distinguish two cases for *while* and *if* statement:

```
switch (P){
case whileStat(_, _):
  println("A while statement");
case ifStat(_, _, _):
  println("An if statement");
}
```

The switch statement (and related statements like *visit*) are even more general than their counterpart in languages like C and Java. Each case is comparable to a *transaction*: when the pattern succeeds and the following statement is executed successfully, all changes to variables made by the statement are *committed* and thus become permanent.

The variables bound by the pattern are always local to the statement associated with the case. When a match fails or when the associated statement fails, a *rollback* to the execution point just before the case takes place and all side-effects are undone. External side-effects like I/O and side-effects in user-defined Java code are not undone. Our motivation for this rather heavy mechanism is that it allows a mostly functional programming style even in cases where the statements associated with a matching pattern lead to failure. In those cases, the side-effects are undone and the next pattern of the switch is tried in exactly the same state as the previous one.

2.8 Visiting

Visiting the elements of a data structure is one of the most common operations in our domain and the visitor design pattern is a solution known to every software engineer. Given a tree-like data structure we want to perform an operation on some (or all) nodes of the tree. The purpose of the visitor design pattern is to decouple the logistics of visiting each node from the actual operation on each node. In Rascal the logistics of visiting is completely automated.

Visiting is achieved by way of visit expressions that resemble the switch statement. A visit expression traverses an arbitrarily complex subject value and applies a number of cases (defined in the same way as cases in a switch statement) to all its subtrees. All the elements of the subject are visited and when one of the cases matches the statements associated with that case are executed. These cases may:

- cause some side effect, i.e., assign a value to local or global variables;
- execute an `insert` statement that replaces the current element;
- execute a `fail` statement that causes the match for the current case to fail (and undoing all side-effects due to the successful match itself and the execution of the statements so far).

The value of a visit expression is the original subject value with all replacements made as dictated by matching cases. The traversal order in a visit expressions can be explicitly defined by the programmer. An example of visiting is given in the next subsection and in Section 3.1, “Colored Trees”.

2.9 Functions

Functions allow the definition of frequently used operations. They have a name and formal parameters. They are explicitly declared and are fully typed. Here is an example of a function that counts the number of assignment statements in a program:

```
int countAssignments(PROGRAM P) {
    int n = 0;
    visit (P) {
        case asgStat(_, _):
            n += 1;
    }
}
```



```

    }
    return n;
}

```

Functions can also be used as values, thus enabling higher-order functions. Consider the following declarations:

```

int double(int x) { return 2 * x; }

int triple(int x) { return 3 * x; }

int f(int x, int (int) multi){ return multi(x); }

```

The functions `double` and `triple` simply multiply their argument with a constant. Function `f` is, however, more interesting. It takes an integer `x` and a function `multi` (with integer argument and integer result) as argument and applies `multi` to its own argument. `f(5, triple)` will hence return 15. Function values can also be created anonymously as illustrated by the following, alternative, manner of writing this same call to `f`:

```

f(5, int (int y){return 3 * y;});

```

Here the second argument of `f` is an anonymous function.

Rascal is a higher-order language in which functions are first-class values. Our motivation to include first-class functions is to provide sufficient mechanisms for writing re-usable analysis and transformation functions.

2.10 Syntax Definition and Parsing

All source code analysis projects need to extract information directly from the source code. There are two main approaches to this:

- *Lexical information*: Use regular expressions to extract useful, but somewhat superficial, flat, information.
- *Structured information*: Use syntax analysis to extract the complete, nested, structure of the source code in the form of a syntax tree.

In Rascal, we reuse the Syntax Definition Formalism (SDF) and its tooling. See <http://www.rascal-mpl.org/> [http://www.meta-environment.org/Meta-Environment/Documentation] for tutorials and manuals for SDF.

SDF modules define grammars and these modules can be imported in a Rascal module. These grammar rules can be applied in writing concrete patterns to match parts of parsed source code. Here is an example of the same while-pattern we saw in Section 2.3, “Pattern Matching”, but now in concrete form:

```

while <Exp> do <Stats> od

```

Importing an SDF module has the following effects:

- All non-terminals (*sorts* in SDF jargon) that are used in the imported grammar are implicitly declared as Rascal types. For each SDF sort S , composite symbols like S^* and $\{S\}$ also become available as types. This makes it possible to handle parse trees and parse tree fragments as fully typed values. They can be assigned to variables, can be stored in larger data structures, can be passed as arguments to functions and can be used in pattern matching.
- For all *start symbols* of the grammar, *parse functions* are implicitly declared that can parse source files according to a specific start symbol.
- Concrete syntax patterns for that specific grammar can be used.
- Concrete syntax constructors can be used that allow the construction of new parse trees.

The following example parses a Java compilation unit from a text file and counts the number of method declarations:

```
module Count
import languages::java::syntax::Java;
import ParseTree;

public int countMethods(loc file){
    int n = 0;
    for(/MethodDeclaration md <- parse(#CompilationUnit,
                                     file))
        n += 1;
    return n;
}
```

First observe that importing the Java grammar has as effect that non-terminals like `MethodDeclaration` and `CompilationUnit` become available as types in the Rascal program.

The implicitly declared function `parse` takes a reified type (`#CompilationUnit`, recall that all non-terminals are implicitly declared as types) and a location as arguments and parses the contents of the location according to the given non-terminal. Next, a match for embedded `MethodDeclarations` is done in the enumerator of the `for` statement. This example ignores many potential error conditions (like error in opening or reading the file, or syntax error during parsing) but does illustrate some of Rascal's syntax and parsing features.

2.11 Rewrite Rules

A *rewrite rule* is a recipe for simplifying values. An example is $(a + b)^2 = a^2 + 2ab + b^2$. A rewrite rule has a pattern as left-hand side (here: $(a + b)^2$) and a

replacement as right-hand side (here: $a^2 + 2ab + b^2$). Given a value and a set of rewrite rules the patterns are tried on every subpart of the value and replacements are made if a match is successful. This is repeated as long as some pattern matches.

Rewrite rules are the only implicit control mechanism in the language and are used to maintain invariants during computations. For example, in a package for symbolic differentiation it is desirable to keep expressions in simplified form in order to avoid intermediate results like `sum(product(num(1), x), product(num(0), y))` that can be simplified to `x`. The following rules achieve this:

```
rule simplify1 product(num(1), Expression e) => e;
rule simplify2 product(Expression e, num(1)) => e;
rule simplify3 product(num(0), Expression e) => num(0);
rule simplify4 product(Expression e, num(0)) => num(0);
rule simplify5 sum(num(0), Expression e)      => e;
rule simplify6 sum(Expression e, num(0))      => e;
```

Whenever a new value of type `Expression` is constructed, these rules are *implicitly* applied to that expression and all its subexpressions. When a pattern at the left-hand side of a rule applies, the matching subexpression is replaced by the right-hand side of the rule. This is repeated as long as any rule can be applied.

Rewrite rules are activated automatically when a value of some specific type is created and one may always assume that values of that type are in simplified form, i.e., all applicable rewrite rules have been applied to it.

Rewrite rules are *Turing complete*, in other words any computable function can be defined using rewrite rules, including functions that do not terminate. The programmer should be aware of this when defining rewrite rules.

2.12 Equation Solving

Many problems can be solved by forms of *constraint solving*. This is a declarative way of programming: specify the constraints that a problem solution should satisfy and how potential solutions can be generated. The actual solution (if any) is found by enumerating possible solutions and testing their compliance with the constraints.

Rascal provides a `solve` statement that helps writing constraint solvers. A typical example is dataflow analysis where the propagation of values through a program can be described by a set of equations. Their solution can be found with the `solve` statement. See Section 5.6, “Dataflow Analysis” for examples.

2.13 Other Features

All language features (including the ones just mentioned) are described in more detail later on in this article. Some features we have not yet mentioned are:

- A Rascal programs consists of a set of modules that are organized in packages.

- A module can import other modules. These can be Rascal modules or SDF modules (as shown above in Section 2.10, “Syntax Definition and Parsing”).
- The visibility of an entity declared in a module can be controlled using a `public/private` modifier.
- A data structures may have annotations that can be explicitly used and modified.
- There is an extensive library for built-in data types, input/output, fact extraction from Java source code, visualization, and more.

2.14 Typechecking and Execution

Rascal has a statically checked type system that prevents type errors and uninitialized variables at runtime. There are no runtime type casts as in Java and there are therefore fewer opportunities for run-time errors. The language provides *higher-order*, *parametric polymorphism*. A type aliasing mechanism allows documenting specific uses of a type. Built-in operators are heavily overloaded. For instance, the operator `+` is used for addition on integers and reals but also for list concatenation, set union etc.

The flow of Rascal program execution is completely explicit. Boolean expressions determine choices that drive the control structures. Rewrite rules form the only exception to the explicit control flow principle. Only local backtracking is provided in the context of boolean expressions and pattern matching; side effects are undone in case of backtracking.

3 Some Simple Examples

The following simple examples will help you to grasp the main features of Rascal quickly. You can also consult the online documentation at <http://www.rascal-mpl.org/> for details of the language or specific operators or functions.

3.1 Colored Trees

Suppose we have binary trees---trees with exactly two children---that have integers as their leaves. Also suppose that our trees can have red and black nodes. Such trees can be defined as follows:

```
module demo::ColoredTrees

data ColoredTree =
    leaf(int N)
  | red(ColoredTree left, ColoredTree right)
  | black(ColoredTree left, ColoredTree right);
```

We can use them as follows:

```
rascal> import demo::ColoredTrees;
```

```
ok
```

```
rascal> rb = red(black(leaf(1), red(leaf(2),leaf(3))),  
         black(leaf(3), leaf(4)));  
ColoredTree: red(black(leaf(1),red(leaf(2),leaf(3))),  
                 black(leaf(3),leaf(4)))
```

Observe that the type of variable `rb` was automatically inferred to be `ColoredTree`.

We define two operations on `ColoredTrees`, one to count the red nodes, and one to sum the values contained in all leaves:

```
// continuing module demo::ColoredTrees  
  
public int cntRed(ColoredTree t){  
    int c = 0;  
    visit(t) {  
        case red(_,_): c = c + 1;❶  
    };  
    return c;  
}  
  
public int addLeaves(ColoredTree t){  
    int c = 0;  
    visit(t) {  
        case leaf(int N): c = c + N;❷  
    };  
    return c;  
}
```

- ❶** Visit all the nodes of the tree and increment the counter `c` for each red node.
- ❷** Visit all nodes of the tree and add the integers in the leaf nodes.

This can be used as follows:

```
rascal> cntRed(rb);  
int: 2  
rascal> addLeaves(rb);  
int: 13
```

A final touch to this example is to introduce green nodes and to replace all red nodes by green ones:

```
// continuing module demo::ColoredTrees  
  
data ColoredTree = green(ColoredTree left,  
                         ColoredTree right);❶
```

```
public ColoredTree makeGreen(ColoredTree t){
    return visit(t) {
        case red(l, r) => green(l, r)      ❷
    };
}
```

- ❶ Extend the ColoredTree data type with a new green constructor.
- ❷ Visit all nodes in the tree and replace red nodes by green ones. Note that the variables `l` and `r` are introduced here without a declaration.

This is used as follows:

```
rascal> makeGreen(rb);
ColoredTree: green(black(leaf(1), green(leaf(2), leaf(3))),
                  black(leaf(3), leaf(4)))
```

This example illustrates the following:

- Abstract data types can be extended as shown by adding the `green` constructor. In large applications this is often encountered as dialects/variants/extensions of a base data type.
- Using `visit` one only has to consider the constructors of interest (e.g., the `leaf` nodes or the `red` nodes) as opposed to all constructors of a datatype. For large data types describing, for instance, programming languages or file formats this can make a large difference in the number of cases that has to be specified.

3.2 Word Replacement

Suppose you are in the publishing business and are responsible for the systematic layout of publications. Authors do not systematically capitalize words in titles---"Word replacement" instead of "Word Replacement"--- and you want to correct this. We solve this problem in two steps. The first step is to define the capitalization of a single word:

```
module demo::WordReplacement
import String;

public str capitalize(str word)
{
    if(/^<letter:[a-z]><rest:.*$>/ := word) ❶
        return toUpperCase(letter) + rest; ❷
    else
        return word; ❸
}
```

- ❶ The function `capitalize` takes a string as input and capitalizes its first character if that is a letter. This is done using a regular expression match that anchors the

match at the beginning (^), expects a single letter and assigns it to the variable `letter` (`letter: [a-z]`) followed by an arbitrary sequence of letters until the end of the string that is assigned to the variable `rest` (`<rest: .*>`).

- ❷ If the regular expression matches we return a new string with the first letter capitalized.
- ❸ Otherwise we return the word unmodified.

The second step is to capitalize all the words in a string. Here are two solutions:

```
// continuing module demo::WordReplacement

public str capAll1(str S)
{
  result = "";
  while (/^<before:\W*><word:\w+><after:.*>/ := S) { ❶
    result += before + capitalize(word);
    S = after;
  }
  return result;
}

public str capAll2(str S)
{
  return visit(S){❷
    case /<word:\w+>/i ❸ => capitalize(word)❹
  };
}
```

- ❶ In the first solution `capAll1` we just loop over all the words in the string and capitalize each word. The variable `result` is used to collect the successive capitalized words. Here we use `\W` to denote non-word characters and `\w` for word characters.
- ❷ In the second solution we use a `visit` expression to visit all the substrings of `S`. Each matching case advances the substring by the length of the pattern it matches and replaces that pattern by another string. If no case matches the next substring is tried.
- ❸ The single case matches a word (note that `\w` matches a word character).
- ❹ When the case matches a word, it is replaced by a capitalized version. The modifier `i` at the end of the regular expressions denotes case-insensitive matching.

We can apply this all as follows:

```
rascal> import demo::WordReplacement;
ok

rascal> capitalize("rascal");
```

```
str: "Rascal"

rascal> capAll1("rascal is great");
str: "Rascal Is Great"
```

This example illustrates that regular expressions are an integral part of the language and can be used in control statements (`if`, `while`), explicit pattern match (`:=`) and as patterns in cases of a `visit`. Their use in the EASY domain is:

- Extracting lexical facts from source code (e.g., identifiers, call statements).
- Parsing DSLs with a strong lexical bias, e.g., markup languages.
- Matching lexical entities in a parse tree, e.g., find all assignments to variables containing the substring "year".

3.3 Template Programming

Many websites and code generators use template-based code generation. They start from a text template that contains embedded variables and code. The template is "executed" by replacing the embedded variables and code by their string value. A language like PHP is popular for this feature. Let's see how we can do this in Rascal. Given a mapping from field names to their type, the task at hand is to generate a Java class that contains those fields and corresponding getters and setters. Given a mapping

```
public map[str, str] fields = (
    "name" : "String",
    "age"  : "Integer",
    "address" : "String"
);
```

we want the call

```
genClass("Person", fields)
```

to produce the following output:

```
public class Person {

    private Integer age;
    public void setAge(Integer age) {
        this.age = age;
    }
    public Integer getAge() {
        return age;
    }
}
```



```

    private String name;
    public void setName(String name) {
        this.name = name;
    }
    public String getName() {
        return name;
    }

    private String address;
    public void setAddress(String address) {
        this.address = address;
    }
    public String getAddress() {
        return address;
    }
}

```

This is achieved by the following definition of `genClass`:

```

module demo::StringTemplate

import String;

public str capitalize(str s) {
    return toUpperCase(substring(s, 0, 1)) +
        substring(s, 1);
}

public str genClass(str name, map[str,str] fields) {
    return "
        public class <name> {
            <for (x <- fields) {
                str t = fields[x];
                str n = capitalize(x);>
                private <t> <x>;
                public void set<n>(<t> <x>) {
                    this.<x> = <x>;
                }
                public <t> get<n>() {
                    return <x>;
                }
            }
        }
    <>>
    "
}

```

This code makes extensive use of Rascal's string interpolation mechanism. All characters inside the string quotes (" and ") are taken literally, except for interpolations that are indicated by angle brackets(< and >). An interpolation is an expression that is evaluated when the string is constructed; the value of the expression replaces the interpolation in the resulting string. Interpolation may contain nested interpolations. The above example contains several examples:

- The interpolation <name > consists of a single variable and will, in this example, be replaced by "Person ". Observe that spaces and layout are significant inside interpolations.
- The interpolation <for (x <- fields){ str t = fields[x]; str n = capitalize(x);> ... <> consists of a complete for statement that will generate all field names, getters and setters. Observe that the statement consists of three parts:
 - The header <for (x <- fields){ str t = fields[x]; str n = capitalize(x);>. Inside the header spaces are not significant.
 - The body private <t> <x> Inside the body spaces are significant. Embedded interpolations may occur inside the body. Observe how these embedded interpolations access the map fields and customize the template for a Java class.
 - The closing brace <}> of the for statement.

This example illustrates how string templates can be used for code generation tasks.

3.4 A Domain-Specific Language for Finite State Machines

Finite State Machines (FSMs) are a universal device in Computer Science and are used to model problems ranging from lexical tokens to concurrent processes. An FSM consists of named states and labeled transitions between states. An example is shown in Figure 1.6, "Example of a Finite State Machine". This example was suggested by G. Hedin at GTTSE09.




Fig. 1.6. Example of a Finite State Machine

This same information can be represented in textual form as follows:

```
finite-state machine
  state S1;
```

```

state S2;
state S3;
trans a: S1 -> S2;
trans b: S2 -> S1;
trans a: S1 -> S3

```

and here is where the idea is born to design a Domain-Specific Language for finite state machines (aptly called FSM). The design of a DSL always proceeds in three steps:

1. **Do domain analysis.** Explore the domain and make an inventory of the relevant concepts and their interactions.
2. **Define syntax.** Design a textual syntax to represent these concepts and interactions.
3. **Define operations.** Define operations on DSL programs. This may be, for example, typechecking, validation, or execution.

We will now apply these steps to the FSM domain.

Do domain analysis. We assume that the FSM domain is sufficiently known. The concepts are states and labeled transitions.

Define syntax. We define a textual syntax for FSMs. This syntax is written in the Syntax Definition Formalism SDF. See <http://www.meta-environment.org/Meta-Environment/Documentation> for tutorials and manuals for SDF. The syntax definition looks as follows:

```

module demo/StateMachine/Syntax

imports basic/Whitespace
imports basic/IdentifierCon

exports
  context-free start-symbols
    FSM

  sorts FSM Decl Trans State IdCon

  context-free syntax
    "state" IdCon                                -> State
    "trans" IdCon ":" IdCon  "->" IdCon          -> Trans
    State                                           -> Decl
    Trans                                           -> Decl
    "finite-state" "machine" {Decl ";" }+        -> FSM

```

Two standard modules for whitespace and identifiers are imported and next a fairly standard grammar for state machines is defined. Observe that SDF syntax rules are

written in (horizontally) reverse order as compared to standard BNF notation, with the nonterminal appearing on the right.

Define Operations. There are various operations one could define on a FSM: executing it for given input tokens, reducing a non-deterministic automaton to a deterministic one, and so on. Here we select a reachability check on FSMs as an example.

We start with the usual imports and define a function `getTransitions` that extracts all transitions from an FSM:

```
module demo::StateMachine::CanReach

import demo::StateMachine::Syntax;
import Relation;
import Map;
import IO;

// Extract from a given FSM all transitions as a relation

public rel[str, str] getTransitions(FSM fsm){
  return
  {<"<from>", "<to>"> |
    /`trans <IdCon a>: <IdCon from> -> <IdCon to>` <- fsm
  };
}
```

The function `getTransitions` illustrates several issues. Given a concrete FSM, a deep pattern match (/) is done searching for `trans` constructs. Since `FSM` is a type that corresponds to a nonterminal, we use a concrete pattern to achieve this: it is enclosed by backquotes (``` and ```) and consists of a grammar rule with embedded variable declarations. For each match three identifiers (`IdCon`) are extracted and assigned to the variables `a`, `from`, respectively, `to`. Next `from` and `to` are converted to a string (using the string interpolations `"<from>"` and `"<to>"`) and finally they are placed in a tuple in the resulting relation. The net effect is that transitions encoded in the syntax tree of `fsm` are collected in a relation for further processing.

Next, we compute all reachable states in the function `canReach`:

```
// continuing module demo::StateMachine::CanReach

public map[str, set[str]] canReach(FSM fsm){
  transitions = getTransitions(fsm);
  return
  ( s: (transitions+)[s] |
    str s <- carrier(transitions)
  );
}
```

Here `str s <- carrier(transitions)` enumerates all elements that occur in the relations that is extracted from `fsm`. A map comprehension is used to construct a map from each state to all states that can be reached it. Transitive closure is denoted by a postfix `+`-operator and `transitions+` is thus the transitive closure of the transition relation and `(transitions+) [s]` gives the image of that closure for a given state; in other words all states that can be reached from it.

Finally, we declare an example FSM (observe that it uses FSM syntax in Rascal code!):

```
// continuing module demo::StateMachine::CanReach

public FSM example =
  `finite-state machine
    state S1;
    state S2;
    state S3;
    trans a: S1 -> S2;
    trans b: S2 -> S1;
    trans a: S1 -> S3`;
```

Testing the above functions gives the following results:

```
rascal> import demo::StateMachine::CanReach;
ok

rascal> getTransitions(example);

rel[str,str]: {<"S1", "S2">, <"S2", "S1">, <"S1", "S3">}
rascal> canReach(example);

map[str: set[str]: ("S1" : {"S1", "S2", "S3"},
                        "S2" : {"S1", "S2", "S3"},
                        "S3" : {})]
```

This example illustrates:

- The use of concrete syntax to define a DSL, e.g., `demo/StateMachine/Syntax`.
- The use of non-terminals of the grammar as types in a Rascal program, e.g., `FSM`.
- The use of quoted DSL fragments in Rascal code, e.g., `example`.
- Extraction of information from a parsed DSL program, e.g., `getTransitions`.
- Representation of this information as relation.
- Description of a reachability check on this relation (e.g., `canReach`).

4 Problem Solving Strategies

Before we study more complicated examples, it is useful to discuss some general problem solving strategies that are relevant in Rascal's application domain. Rascal and supporting libraries and tools were specifically designed to support these strategies.

To appreciate these general strategies, it is good to keep some specific problem areas in mind:

- **Documentation generation:** extract facts from source code and use them to generate textual documentation. A typical example is generating web-based documentation for legacy languages like Cobol and PL/I.
- **Metrics calculation:** extract facts from source code (and possibly other sources like test runs) and use them to calculate code metrics. Examples are cohesion and coupling of modules and test coverage.
- **Model extraction:** extract facts from source code and use them to construct an abstract model of the source code. An example is extracting lock and unlock calls from source code and building an automaton that guarantees that lock/unlock occurs in pairs along every control flow path.
- **Model-based code generation:** given a high-level model of a software system, described in UML or some other modelling language, transform this model into executable code. UML-to-Java code generation falls in this category.
- **Source-to-source transformation:** given certain objectives like removing deprecated language features, upgrading to newer APIs and the like, perform large-scale, fully automated, source code transformation.
- **Interactive refactoring:** given known "code smells" allow a user to interactively indicate how these smells should be removed. The refactoring features in Eclipse and Visual Studio are examples.

With these examples in mind, we can study the overall problem solving workflow as shown in Figure 1.7, "General 3-Phased Problem Solving Workflow". It consists of three optional phases:

- If **extraction** is needed to solve the problem, then define the extraction phase, see Section 4.1, "Defining Extraction".
- If **analysis** is needed, then define the analysis phase, see Section 4.2, "Defining Analysis".
- If **synthesis** is needed, then define the synthesis phase, see Section 4.3, "Defining Synthesis".




Fig. 1.7. General 3-Phased Problem Solving Workflow

Each phase is subject to a validation and improvement workflow as shown in Figure 1.8, “Validation and Improvement Workflow”. Each individual phase as well as the combination of phases may introduce errors and has thus to be carefully validated. In combination with the detailed strategies for each phase, this forms a complete approach for problem solving and validation using Rascal.




Fig. 1.8. Validation and Improvement Workflow

A major question in every problem solving situation is how to determine the requirements for each phase of the solution. For instance, how do we know what to extract from the source code if we do not know what the desired end results of the project are? The standard solution is to use a workflow for requirements gathering that

is the inverse of the phases needed to solve the complete problem. This is shown in Figure 1.9, “Requirements Workflow” and amounts to the phases:

- **Requirements of the synthesis phase.** This amounts to making an inventory of the desired results of the whole project and may include generated source code, abstract models, or visualizations.
- **Requirements of the analysis phase.** Once these results of the synthesis phase are known, it is possible to list the analysis results that are needed to synthesize desired results. Possible results of the analysis phase include type information, structural information of the original source.
- **Requirements of the extraction phase.** As a last step, one can make an inventory of the facts that have to be extracted to form the starting point for the analysis phase. Typical facts include method calls, inheritance relations, control flow graphs, usage patterns of specific library functions or language constructs.

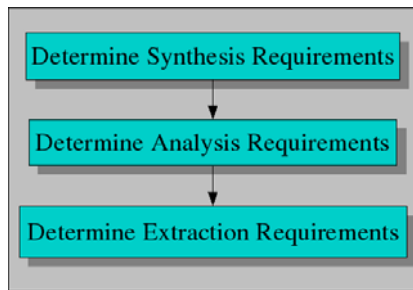


Fig. 1.9. Requirements Workflow

You will have to identify requirements for each phase when you apply them to a specific example from the list given earlier.

When these requirements have been established, it becomes much easier to actually carry out the project using the three phases of Figure 1.7, “General 3-Phased Problem Solving Workflow”.

4.1 Defining Extraction

How can we extract facts from the *System under Investigation* (SUI) that we are interested in? The extraction workflow is shown in Figure 1.10, “Extraction Workflow” and consists of the following steps:

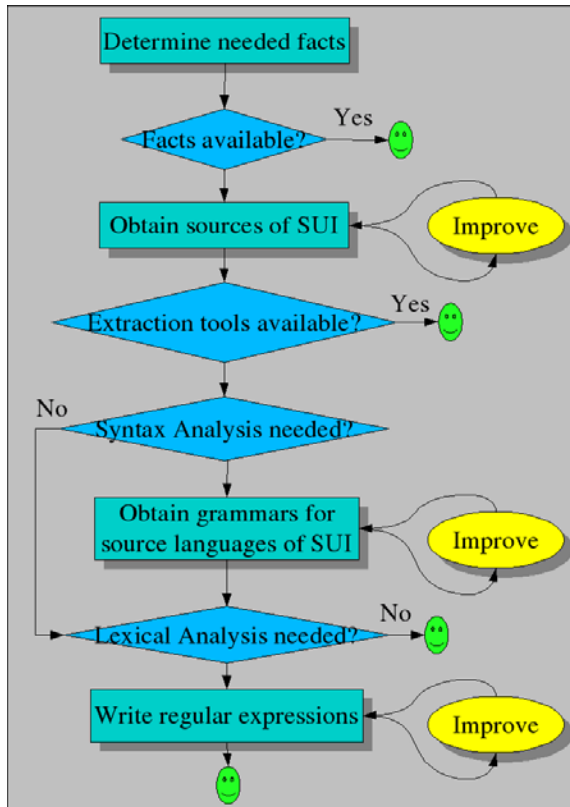


Fig. 1.10. Extraction Workflow

- First and foremost we have to determine which facts we need. This sounds trivial, but it is not. A common approach is to use look-ahead and to sketch the queries that are likely to be used in the analysis phase and to determine which facts are needed for them. Start with extracting these facts and refine the extraction phase when the analysis phase is completely defined.
- If relevant facts are already available (and they are reliable!) then we are done. This may happen when you are working on a system that has already been analyzed by others.
- Otherwise you need the source code of the SUI. This requires:

- Checking that all sources are available (and can be compiled by the host system on which they are usually compiled and executed). Due to missing or unreliable configuration management on the original system this may be a labour-intensive step that requires many iterations.
- Determining in which languages the sources are written. In larger systems it is common that three or more different languages are being used.
- If there are reliable third-party extraction tools available for this language mix, then we only have to apply them and we are done. Here again, validation is needed that the extracted facts are as expected.
- The extraction may require syntax analysis. This is the case when more structural properties of the source code are needed such as the flow-of-control, nesting of declarations, and the like. There two approaches here:
 - Use a third-party parser, convert the source code to parse trees and do the further processing of these parse trees in Rascal. The advantage is that the parser can be re-used, the disadvantage is that data conversion is needed to adapt the generated parse tree to Rascal. Validate that the parser indeed accepts the language the SUI is written in, since you will not be the first who has been bitten by the language dialect monster when it turns out that the SUI uses a local variant that slightly deviates from a mainstream language.
 - Use an existing SDF definition of the source language or write your own definition. In both cases you can profit from Rascal's seamless integration with SDF. Be aware, however, that writing a grammar for a non-trivial language is a major undertaking and may require weeks to months of work. Whatever approach you choose, validate the resulting grammar.
- The extraction phase may only require lexical analysis. This happens when more superficial, textual, facts have to be extracted like procedure calls, counts of certain statements and the like. Use Rascal's full regular expression facilities to do the lexical analysis.

It may happen that the facts extracted from the source code are *wrong*. Typical error classes are:

- Extracted facts are *wrong*: the extracted facts incorrectly state that procedure P calls procedure Q but this is contradicted by a source code inspection. This may happen when the fact extractor uses a conservative approximation when precise information is not statically available. In the C language, when procedure P performs an indirect call via a pointer variable, the approximation may be that P calls all procedures in the program.
- Extracted facts are *incomplete*: for example, the inheritance between certain classes in Java code is missing.

The strategy to validate extracted facts differs per case but here are three possibilities:

- Post process the extracted facts (using Rascal, of course) to obtain trivial facts about the source code such as total lines of source code and number of procedures, classes, interfaces and the like. Next validate these trivial facts with tools like `wc` (word and line count), `grep` (regular expression matching) and others.
- Do a manual fact extraction on a small subset of the code and compare this with the automatically extracted facts.
- Use another tool on the same source and compare results whenever possible. A typical example is a comparison of a call relation extracted with different tools.

The Rascal features that are most frequently used for extraction are:

- Regular expression patterns to extract textual facts from source code.
- Syntax definitions and concrete patterns to match syntactic structures in source code.
- Pattern matching (used in many Rascal statements).
- Visits to traverse syntax trees and to locally extract information.
- The repertoire of built-in data types (like lists, maps, sets and relations) to represent the extracted facts.

A large diversity of problems can be encountered in the extraction phase. Rascal tries to provide a uniform framework with the right tools to solve them.

4.2 Defining Analysis

The analysis workflow is shown in Figure 1.11, “Analysis Workflow” and consists of two steps:

- Determine the results that are needed for the synthesis phase.
- Write the Rascal code to perform the analysis. This may amount to:
 - Reordering extracted facts to make them more suitable for the synthesis phase.
 - Enriching extracted facts. Examples are computing transitive closures of extracted facts (e.g., A may call B in one or more calls), or performing data reduction by abstracting away details (i.e., reducing a program to a finite automaton).
 - Combining enriched, extracted, facts to create new facts.

As before, validate, validate and validate the results of analysis. Essentially the same approach can be used as for validating the facts. Manual checking of answers on random

samples of the SUI may be mandatory. It also happens frequently that answers inspire new queries that lead to new answers, and so on.

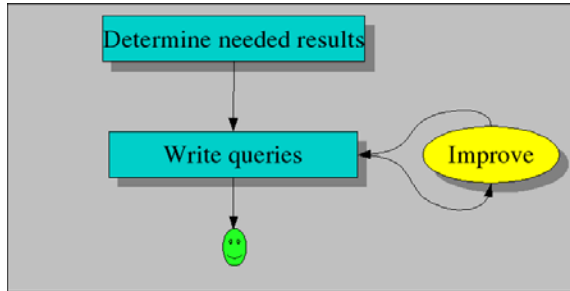


Fig. 1.11. Analysis Workflow

The Rascal features that are frequently used for analysis are:

- List, set and map comprehensions.
- The built-in operators and library functions, in particular for lists, maps, sets and relations.
- Pattern matching (used in many Rascal statements).
- Visits and switches to further process extracted facts.
- The solve statement for constraint solving.
- Rewrite rules to simplify results and to enforce constraints.

4.3 Defining Synthesis

Results are synthesized as shown in Figure 1.12, “Synthesis Workflow”. This consists of the following steps:

- Determine the results of the synthesis phase. A wide range of results is possible including:
 - Generated source code.
 - Generated abstract representations, like finite automata or other formal models that capture properties of the SUI.
 - Generated data for visualizations that will be used by visualization tools.
- If source code is to be generated, there are various options.
 - Print strings that are customized using string interpolation.

- First generate the desired output in the form of an abstract syntax tree and then convert this tree to a string (perhaps using forms of pretty printing).
- Generate the desired output directly as syntactically correct structured text. This is achieved by using a grammar of the target source language and composing syntactically correct program fragments during code generation. The grammar defines the structure of parse trees for the target language and guarantees the construction of type correct parse tree fragments. This approach hence guarantees the generation of syntactically correct source code as opposed to code generation using print statements or string templates.
- If other output is needed (e.g., an automaton or other formal structure) write data declarations to represent that output.
- Finally, write functions and rewrite rules that generate the desired results.

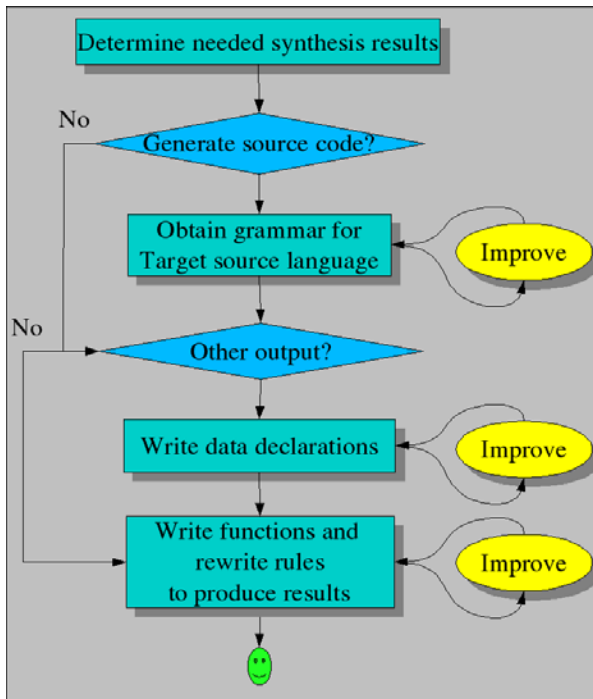


Fig. 1.12. Synthesis Workflow

The Rascal features that are frequently used for synthesis are:

- Syntax definitions or data declarations to define output formats.
- Pattern matching (used in many Rascal statements).

- Visits of data structures and on-the-fly code generation.
- Rewrite rules.

5 Larger Examples

Now we will have a closer look at some larger applications of Rascal. We start with a call graph analysis in Section 5.1, “Call Graph Analysis” and then continue with the analysis of the component structure of an application in Section 5.2, “Analyzing the Component Structure of an Application” and of Java systems in Section 5.3, “Analyzing the Structure of Java Systems”. Next we move on to the detection of uninitialized variables in Section 5.4, “Finding Uninitialized and Unused Variables in a Program”. As an example of computing code metrics, we describe the calculation of McCabe’s cyclomatic complexity in Section 5.5, “McCabe Cyclomatic Complexity”. Several examples of dataflow analysis follow in Section 5.6, “Dataflow Analysis”. A description of program slicing concludes the chapter, see Section 5.7, “Program Slicing”.

5.1 Call Graph Analysis

Suppose a mystery box ends up on your desk. When you open it, it contains a huge software system with several questions attached to it:

- How many procedure calls occur in this system?
- How many procedures does it contains?
- What are the entry points for this system, i.e., procedures that call others but are not called themselves?
- What are the leaves of this application, i.e., procedures that are called but do not make any calls themselves?
- Which procedures call each other indirectly?
- Which procedures are called directly or indirectly from each entry point?
- Which procedures are called from all entry points?

Let’s see how these questions can be answered using Rascal.

5.1.1 Preparations

To illustrate this process consider the workflow in Figure 1.13, “Workflow for analyzing mystery box”. First we have to extract the calls from the source code. Rascal is very good at this, but to simplify this example we assume that this call graph has already been extracted. Also keep in mind that a real call graph of a real application will contain thousands and thousands of calls. Drawing it in the way we do later on in Figure 1.14, “Graphical representation of the `calls` relation” makes no sense since we get a uniformly black picture due to all the call dependencies. After the extraction phase,

we try to understand the extracted facts by writing queries to explore their properties. For instance, we may want to know *how many calls* there are, or *how many procedures*. We may also want to enrich these facts, for instance, by computing who calls who in more than one step. Finally, we produce a simple textual report giving answers to the questions we are interested in.




Fig. 1.13. Workflow for analyzing mystery box

Now consider the call graph shown in Figure 1.14, “Graphical representation of the `calls` relation”. This section is intended to give you a first impression what can be done with Rascal.




Fig. 1.14. Graphical representation of the `calls` relation

Rascal supports basic data types like integers and strings which are sufficient to formulate and answer the questions at hand. However, we can gain readability by introducing separately named types for the items we are describing. First, we introduce therefore a new type `proc` (an alias for strings) to denote procedures:

```
rascal> alias proc = str;
ok
```

Suppose that the following facts have been extracted from the source code and are represented by the relation `Calls`:

```
rascal> rel[proc, proc] Calls =
    { <"a", "b">, <"b", "c">, <"b", "d">, <"d", "c">,
      <"d", "e">, <"f", "e">, <"f", "g">, <"g", "e">
    };
rel[proc,proc]: { <"a", "b">, <"b", "c">, <"b", "d">,
                  <"d", "c">, <"d", "e">, <"f", "e">,
                  <"f", "g">, <"g", "e">}
```

This concludes the preparatory steps and now we move on to answer the questions.

5.1.2 How Many Procedure Calls Occur in This System?

To determine the numbers of calls, we simply determine the number of tuples in the `Calls` relation, as follows. First, we need the `Relation` library so we import it:

```
rascal> import Relation;
ok
```

next we describe a new variable and calculate the number of tuples:

```
rascal> nCalls = size(Calls);
int: 8
```

The library function `size` determines the number of elements in a set or relation. In this example, `nCalls` will get the value 8.

5.1.3 How Many Procedures Are Contained in It?

We get the number of procedures by determining which names occur in (the first or second component of) the tuples in the relation `Calls` and then determining the number of names:

```
rascal> procs = carrier(Calls);
set[proc]: {"a", "b", "c", "d", "e", "f", "g"}

rascal> nprocs = size(procs);
int: 7
```

The built-in function `carrier` determines all the values that occur in the tuples of a relation. In this case, `procs` will get the value `{"a", "b", "c", "d", "e", "f", "g"}` and `nprocs` will thus get value 7. A more concise way of expressing this would be to combine both steps:

```
rascal> nprocs = size(carrier(Calls));
int: 7
```


5.1.4 What Are the Entry Points for This System?

The next step in the analysis is to determine which *entry points* this application has, i.e., procedures which call others but are not called themselves. Entry points are useful since they define the external interface of a system and may also be used as guidance to split a system in parts. The `top` of a relation contains those left-hand sides of tuples in a relation that do not occur in any right-hand side. When a relation is viewed as a graph, its `top` corresponds to the root nodes of that graph. Similarly, the `bottom` of a relation corresponds to the leaf nodes of the graph. Using this knowledge, the entry points can be computed by determining the `top` of the `Calls` relation:

```
rascal> import Graph;
ok

rascal> entryPoints = top(Calls);
set[proc]: {"a", "f"}
```

In this case, `entryPoints` is equal to `{"a", "f"}`. In other words, procedures "a" and "f" are the entry points of this application.

5.1.5 What Are the Leaves of This Application?

In a similar spirit, we can determine the *leaves* of this application, i.e., procedures that are being called but do not make any calls themselves:

```
rascal> bottomCalls = bottom(Calls);
set[proc]: {"c", "e"}
```

In this case, `bottomCalls` is equal to `{"c", "e"}`.

5.1.6 Which Procedures Call Each Other Indirectly?

We can also determine the *indirect calls* between procedures, by taking the transitive closure of the `Calls` relation, written as `Calls+`. Observe that the transitive closure will contain both the direct and the indirect calls.

```
rascal> closureCalls = Calls+;
rel[proc, proc]: {<"a", "b">, <"b", "c">, <"b", "d">,
                  <"d", "e">, <"f", "e">,
                  <"f", "g">, <"g", "e">, <"a", "c">,
                  <"a", "d">, <"b", "e">, <"a", "e">}
```

It is easy to get rid of all the direct calls in the above result, by subtracting (using the set difference operator `-`) the original call relation from it:

```
rascal> indirectClosureCalls = closureCalls - Calls;
rel[proc, proc]: {<"a", "c">, <"a", "d">,
                  <"b", "e">, <"a", "e">}
```

5.1.7 Which Procedures Are Called Directly or Indirectly from Each Entry Point?

We now know the entry points for this application ("a" and "f") and the indirect call relations. Combining this information, we can determine which procedures are called from each entry point. This is done by indexing `closureCalls` with an appropriate procedure name. The index operator yields all right-hand sides of tuples that have a given value as left-hand side. This gives the following:

```
rascal> calledFromA = closureCalls["a"];
set[proc]: {"b", "c", "d", "e"}
```

and

```
rascal> calledFromF = closureCalls["f"];
set[proc]: {"e", "g"}
```

5.1.8 Which procedures are called from all entry points?

Finally, we can determine which procedures are called from both entry points by taking the intersection (&) of the two sets `calledFromA` and `calledFromF`:

```
rascal> commonProcs = calledFromA & calledFromF;
set[proc]: {"e"}
```

In other words, the procedures called from both entry points are mostly disjoint except for the common procedure "e".

5.1.9 Wrap-Up

These findings can be verified by inspecting a graph view of the calls relation as shown in Figure 1.14, "Graphical representation of the `calls` relation". Such a visual inspection does *not* scale very well to large graphs and this makes the above form of analysis particularly suited for studying large systems.

This example illustrates the following:

- Call dependencies can easily be represented as relations.
- By design, Rascal contains the full repertoire of operators and library functions for manipulating relations.
- This style of programming enables flexible exploration of given data.

5.2 Analyzing the Component Structure of an Application

A frequently occurring problem is that we know the call relation of a system but that we want to understand it at the component level rather than at the procedure level. If it is known to which component each procedure belongs, it is possible to *lift* the call relation

to the component level as proposed in [Kri99]. Actual lifting amounts to translating each call between procedures to a call between components. This is described in the following module:

```
module demo::Lift

alias proc = str;
alias comp = str;

public rel[comp,comp] lift(rel[proc,proc] aCalls,
                           rel[proc,comp] aPartOf){
  return
    { <C1, C2> | <proc P1, proc P2> <- aCalls,
                <comp C1, comp C2> <- aPartOf[P1] *
                                     aPartOf[P2]
    };
}
```

For each pair $\langle P1, P2 \rangle$ in the Calls relation we compose the corresponding parts $aPartOf[P1]$ and $aPartOf[P2]$ (each yielding a set of components) into a new relation of calls between components. This relation is added pair by pair to the result.

Let's now apply this. First import the above module, and define a call relation and a partof relation:

```
rascal> import demo::Lift;
ok

rascal> Calls = {<"main", "a">, <"main", "b">, <"a", "b">,
                 <"a", "c">, <"a", "d">, <"b", "d">
                 };
rel[str,str] : {<"main", "a">, <"main", "b">, <"a", "b">,
                <"a", "c">, <"a", "d">, <"b", "d">
                }

rascal> Components = {"Appl", "DB", "Lib"};
set[str] : {"Appl", "DB", "Lib"}

rascal> PartOf = {<"main", "Appl">, <"a", "Appl">,
                  <"b", "DB">, <"c", "Lib">,
                  <"d", "Lib">};
rel[str,str] : {<"main", "Appl">, <"a", "Appl">,
                <"b", "DB">, <"c", "Lib">,
                <"d", "Lib">}
```

The lifted call relation between components is now obtained by:

```
rascal> ComponentCalls = lift(Calls, PartOf);
rel[str,str] : {<"DB", "Lib">, <"Appl", "Lib">,
               <"Appl", "DB">, <"Appl", "Appl">}
```

The relevant relations for this example are shown in Figure 1.15, “(a) **Calls**; (b) **PartOf**; (c) **ComponentCalls**.”.




Fig. 1.15. (a) **Calls**; (b) **PartOf**; (c) **ComponentCalls**.

This example illustrates the following:

- A comprehension to build a relation, e.g., as done in `lift`.
- Relation projection, e.g., `aPartOf[P1]`.
- Set product, e.g., `aPartOf[P1] * aPartOf[P2]`.

5.3 Analyzing the Structure of Java Systems

Now we consider the analysis of Java systems (inspired by [BNL05]). Suppose that the type class is defined as follows

```
alias class = str;
```

and that the following relations are available about a Java application:

- `rel[class,class] CALL`: If $\langle C_1, C_2 \rangle$ is an element of `CALL`, then some method of C_2 is called from C_1 .
- `rel[class,class] INHERITANCE`: If $\langle C_1, C_2 \rangle$ is an element of `INHERITANCE`, then class C_1 either extends class C_2 or C_1 implements interface C_2 .
- `rel[class,class] CONTAINMENT`: If $\langle C_1, C_2 \rangle$ is an element of `CONTAINMENT`, then one of the fields of class C_1 is of type C_2 .

To make this more explicit, consider the class `LocatorHandle` from the `JHotDraw` application (version 5.2) as shown here:

```
package CH.ifa.draw.standard;
```

```

import java.awt.Point;
import CH.ifa.draw.framework.*;
/**
 * A LocatorHandle implements a Handle by delegating the
 * location requests to a Locator object.
 */
public class LocatorHandle extends AbstractHandle {
    private Locator      fLocator;
    /**
     * Initializes the LocatorHandle with the
     * given Locator.
     */
    public LocatorHandle(Figure owner, Locator l) {
        super(owner);
        fLocator = l;
    }
    /**
     * Locates the handle on the figure by forwarding
     * the request to its figure.
     */
    public Point locate() {
        return fLocator.locate(owner());
    }
}

```

It leads to the addition to the above relations of the following tuples:

- To CALL the pairs <"LocatorHandle", "AbstractHandle"> and <"LocatorHandle", "Locator"> will be added.
- To INHERITANCE the pair <"LocatorHandle", "AbstractHandle"> will be added.
- To CONTAINMENT the pair <"LocatorHandle", "Locator"> will be added.

Cyclic structures in object-oriented systems makes understanding hard. Therefore it is interesting to spot classes that occur as part of a cyclic dependency. Here we determine cyclic uses of classes that include calls, inheritance and containment. This is achieved as follows:

```

rel[class,class] USE = CALL + CONTAINMENT + INHERITANCE;
set[str] ClassesInCycle =
    {C1 | <class C1, class C2> <- USE+, C1 == C2};

```

First, we define the USE relation as the union of the three available relations CALL, CONTAINMENT and INHERITANCE. Next, we consider all pairs < C_1 , C_2 > in the

transitive closure of the USE relation such that C_1 and C_2 are equal. Those are precisely the cases of a class with a cyclic dependency on itself. Probably, we want to know not only which classes occur in a cyclic dependency, but we also which classes are involved in such a cycle. In other words, we want to associate with each class a set of classes that are responsible for the cyclic dependency. This can be done as follows.

```
rel[class,class] USE = CALL + CONTAINMENT + INHERITANCE;
set[class] CLASSES = carrier(USE);
rel[class,class] USETRANS = USE+;
rel[class,set[class]] ClassCycles =
    {<C, USETRANS[C]> | class C <- CLASSES,
                        <C, C> in USETRANS };
```

First, we introduce two new shorthands: CLASSES and USETRANS. Next, we consider all classes C with a cyclic dependency and add the pair $\langle C, \text{USETRANS}[C] \rangle$ to the relation ClassCycles. Note that $\text{USETRANS}[C]$ is the right image of the relation USETRANS for element C , i.e., all classes that can be called transitively from class C .

This example illustrates the following:

- Representation of facts related to declarations in Java programs.
- Manipulation of these facts using comprehensions and relational operators.

One of the Rascal libraries gives access to the Eclipse JDT and provides information about declarations and types of Java programs in relational form. This is briefly illustrated in Section 5.8, “Visualizing Extracted Information”.

5.4 Finding Uninitialized and Unused Variables in a Program

Consider the following program in the toy language Pico: (This is an extended version of the example presented earlier in [Kli03].)

```
[ 1] begin declare x : natural, y : natural,
[ 2]           z : natural, p : natural;
[ 3]   x := 3;
[ 4]   p := 4;
[ 5]   if q then
[ 6]       z := y + x
[ 7]   else
[ 8]       x := 4
[ 9]   fi;
[10]   y := z
[11] end
```

Inspection of this program shows that some of the variables are being used before they have been initialized. The variables in question are q (line 5), y (line 6), and z (line

10). It is also clear that variable p is initialized (line 4), but is never used. How can we automate these kinds of analysis? Recall from Section 1.1, “The EASY Paradigm” that we follow the Extract-Analyze-SYnthesize paradigm to approach such a problem. The first step is to determine which elementary facts we need about the program. For this and many other kinds of program analysis, we need at least the following:

- The *control flow graph* of the program. We represent it by a graph $PRED$ (for predecessor) which relates each statement with its predecessors.
- The *definitions* of each variable, i.e., the program statements where a value is assigned to the variable. It is represented by the relation $DEFS$.
- The *uses* of each variable, i.e., the program statements where the value of the variable is used. It is represented by the relation $USES$.

In this example, we will use line numbers to identify the statements in the program. Assuming that there is a tool to extract the above information from a program text, we get the following for the above example:

```
module demo::Uninit
import Relation;
import Graph;

alias expr = int;
alias varname = str;

public expr ROOT = 1;

public graph[expr] PRED = { <1,3>, <3,4>, <4,5>, <5,6>,
                           <5,8>, <6,10>, <8,10> };

public rel[varname,expr] DEFS = { <"x", 3>, <"p", 4>,
                                  <"z", 6>, <"x", 8>,
                                  <"y", 10> };

public rel[varname, expr] USES = { <"q", 5>, <"y", 6>,
                                   <"x", 6>, <"z", 10> };
```

This concludes the extraction phase. Next, we have to enrich these basic facts to obtain the initialized variables in the program. So, when is a variable V in some statement S initialized? If we execute the program (starting in $ROOT$), there may be several possible execution paths that can reach statement S . All is well if *all* these execution path contain a definition of V . However, if one or more of these path do *not* contain a definition of V , then V may be uninitialized in statement S . This can be formalized as follows:

```
// module demo::Unit continued
public rel[varname,expr] UNINIT =
```

```
{ <V,E> | <varname V, expr E> <- USES,
      E in reachX(PRED, {ROOT}, DEFS[V])
};
```

We analyze this definition in detail:

- `<varname V, expr E> : USES` enumerates all tuples in the `USES` relation. In other words, we consider the use of each variable in turn.
- `E in reachX(PRED, {ROOT}, DEFS[V])` is a test that determines whether expression *E* is reachable from the `ROOT` without encountering a definition of variable *V*.
 - `PRED` is the relation for which the reachability has to be determined.
 - `{ROOT}` represents the initial set of nodes from which all path should start.
 - `DEFS[V]` yields the set of all statements in which a definition of variable *V* occurs. These nodes form the exclusion set for `reachX`: no path will be extended beyond an element in this set.
- The result of `reachX(PRED, {ROOT}, DEFS[V])` is a set that contains all nodes that are reachable from the `ROOT`.
- Finally, `E in reachX(PRED, {ROOT}, DEFS[V])` tests whether expression *E* can be reached from the `ROOT`.
- The net effect is that `UNINIT` will only contain pairs that satisfy the test just described.

When we execute the resulting Rascal code (i.e., the declarations of `ROOT`, `PRED`, `DEFS`, `USES` and `UNINIT`), we get as value for `UNINIT`:

```
rascal> import demo::Uninit;
ok

rascal> UNINIT;
rel[varname,expr]: {"q", 5>, "y", 6>, "z", 10>}
```

and this is in concordance with the informal analysis given at the beginning of this example.

As a bonus, we can also determine the *unused* variables in a program, i.e., variables that are defined but are used nowhere. This is done as follows:

```
// module demo::Unit continued

public set[varname] UNUSED = domain(DEFS) - domain(USES);
```


Taking the domain of the relations `DEFS` and `USES` yields the variables that are defined, respectively, used in the program. The difference of these two sets yields the unused variables, in this case `{ "p" }`.

This example illustrates the following:

- Representation of the control flow graph as a relation.
- Reachability computations on that flow graph.

5.5 McCabe Cyclomatic Complexity

The *cyclomatic complexity* of a program is defined as $e - n + 2$, where e and n are the number of edges and nodes in the control flow graph, respectively. It was proposed by McCabe [McC76] as a measure of program complexity. Experiments have shown that programs with a higher cyclomatic complexity are more difficult to understand and test and have more errors. It is generally accepted that a program, module or procedure with a cyclomatic complexity larger than 15 is *too complex*. Essentially, cyclomatic complexity measures the number of decision points in a program. Given a control flow in the form of a predecessor graph `Graph[&T] PRED` between elements of arbitrary type `&T`, the cyclomatic complexity can be computed in Rascal as follows:

```
module demo::McCabe
import Graph;

public int cyclomaticComplexity(Graph[&T] PRED) {
    return size(PRED) - size(carrier(PRED)) + 2;
}
```

The number of edges e is equal to the number of tuples in `PRED`. The number of nodes n is equal to the number of elements in the carrier of `PRED`, i.e., all elements that occur in a tuple in `PRED`.

This example illustrates that metrics on a (control flow) graph can be easily expressed.

5.6 Dataflow Analysis

Dataflow analysis is a program analysis technique that forms the basis for many compiler optimizations. It is described in any text book on compiler construction, e.g. [ASU86]. The goal of dataflow analysis is to determine the effect of statements on their surroundings. Typical examples are:

- Dominators (Section 5.6.1, “Dominator”): which nodes in the flow dominate the execution of other nodes?
- Reaching definitions (Section 5.6.2, “Reaching Definitions”): which definitions of variables are still valid at each statement?

- Live variables (Section 5.6.3, “Live Variables”): of which variables will the values be used by successors of a statement?
- Available expressions: which expressions are computed along each path from the start of the program to the current statement?

5.6.1 Dominators

A node d of a flow graph *dominates* a node n , if every path from the initial node of the flow graph to n goes through d [ASU86] (Section 10.4). Dominators play a role in the analysis of conditional statements and loops. The function `dominators` that computes the dominators for a given flow graph `PRED` and an entry node `ROOT` is defined as follows:

```
module demo::Dominators
import Set;
import Relation;
import Graph;

public rel[&T, set[&T]] dominators(rel[&T,&T] PRED,
                                   &T ROOT)
{
    set[&T] VERTICES = carrier(PRED);
    return { <V, (VERTICES - {V, ROOT}) -
              reachX(PRED, {ROOT}, {V})>
            | &T V <- VERTICES
            };
}
```

First, the auxiliary set `VERTICES` (all the statements) is computed. The relation `DOMINATES` consists of all pairs $\langle S, \{S_1, \dots, S_n\} \rangle$ such that

- S_i is not an initial node or equal to S .
- S_i cannot be reached from the initial node without going through S .

First import the above module and consider the sample flow graph `PRED`:

```
rascal> import demo::Dominators;
ok

rascal> rel[int,int] PRED = {
<1,2>, <1,3>,
<2,3>,
<3,4>,
<4,3>,<4,5>, <4,6>,
<5,7>,
```

```

<6,7>,
<7,4>,<7,8>,
<8,9>,<8,10>,<8,3>,
<9,1>,
<10,7>
};

rel[int,int]: { <1,2>, <1,3>, ...

```

It is illustrated in Figure 1.16, “Flow graph”




Fig. 1.16. Flow graph

The result of applying dominators to it is as follows:

```

rascal> dominators(PRED);
rel[int,int]: {<1, {2, 3, 4, 5, 6, 7, 8, 9, 10}>,
<2, {}>,
<3, {4, 5, 6, 7, 8, 9, 10}>,
<4, {5, 6, 7, 8, 9, 10}>,
<5, {}>,
<6, {}>,
<7, {8, 9, 10}>,
<8, {9, 10}>,
<9, {}>,
<10, {}>}}

```

The resulting *dominator tree* is shown in Figure 1.17, “Dominator tree”. The dominator tree has the initial node as root and each node d in the tree only dominates its descendants in the tree.




Fig. 1.17. Dominator tree




Fig. 1.18. Flow graph for various dataflow problems

5.6.2 Reaching Definitions

We illustrate the calculation of reaching definitions using the example in Figure 1.18, “Flow graph for various dataflow problems” which was inspired by [ASU86] (Example 10.15).

We introduce the notions *definition* and *use* to represent information about the program. The former describes that a certain statement defines some variable and the latter that a statement uses some variable. They are defined as follows:

```

module demo::ReachingDefs

import Relation;
import Graph;
import IO;

```

```

public alias stat = int;
public alias var = str;
public alias def  = tuple[stat, var];
public alias use  = tuple[stat,var];

public rel[stat,def] definition(rel[stat,var] DEFS){
  return {<S,<S,V>> | <stat S, var V> <- DEFS};
}

public rel[stat,def] use(rel[stat, var] USES){
  return {<S, <S, V>> | <stat S, var V> <- USES};
}

```

Let's use the following values to represent our example:

```

rascal> rel[stat,stat] PRED = { <1,2>, <2,3>, <3,4>,
                                <4,5>, <5,6>, <5,7>,
                                <6,7>, <7,4> };
rel[stat,stat]: { <1,2>, <2,3>, ...

rascal> rel[stat, var] DEFS = { <1, "i">, <2, "j">,
                                <3, "a">, <4, "i">,
                                <5, "j">, <6, "a">,
                                <7, "i"> };
rel[stat,var]: { <1, "i">, <2, "j">, ...

rascal> rel[stat,var] USES = { <1, "m">, <2, "n">,
                                <3, "u1">, <4, "i">,
                                <5, "j">, <6, "u2">,
                                <7, "u3"> };
rel[stat,var]: { <1, "m">, <2, "n">, ...

```

The functions `definition` and `use` have the following effect on our sample data:

```

rascal> definition(DEFS);
rel[stat,def]: { <1, <1, "i">>, <2, <2, "j">>,
                 <3, <3, "a">>, <4, <4, "i">>,
                 <5, <5, "j">>, <6, <6, "a">>,
                 <7, <7, "i">> }

rascal> use(USES);
rel[stat,def]: { <1, <1, "m">>, <2, <2, "n">>,
                 <3, <3, "u1">>, <4, <4, "i">>,
                 <5, <5, "j">>, <6, <6, "u2">>,
                 <7, <7, "u3">> }

```

Now we are ready to define an important new relation `KILL`. `KILL` defines which variable definitions are undone (killed) at each statement and is defined by the following function `kill`:

```
// continuing module demo::ReachingDefs

public rel[stat,def] kill(rel[stat,var] DEFS) {
    return {<S1, <S2, V>> | <stat S1, var V> <- DEFS,
                                   <stat S2, V> <- DEFS,
                                   S1 != S2};
}
```

In this definition, all variable definitions are compared with each other, and for each variable definition all *other* definitions of the same variable are placed in its kill set. In the example, `KILL` gets the value

```
rascal> kill(DEFS);
rel[stat,def]:
{ <1, <4, "i">>, <1, <7, "i">>, <2, <5, "j">>,
  <3, <6, "a">>, <4, <1, "i">>, <4, <7, "i">>,
  <5, <2, "j">>, <6, <3, "a">>, <7, <1, "i">>,
  <7, <4, "i">>
}
```

and, for instance, the definition of variable `i` in statement 1 kills the definitions of `i` in statements 4 and 7.

After these preparations, we are ready to formulate the reaching definitions problem in terms of two relations `IN` and `OUT`. `IN` captures all the variable definitions that are valid at the entry of each statement and `OUT` captures the definitions that are still valid after execution of each statement. Intuitively, for each statement `S`, `IN[S]` is equal to the union of the `OUT` of all the predecessors of `S`. `OUT[S]`, on the other hand, is equal to the definitions generated by `S` to which we add `IN[S]` minus the definitions that are killed in `S`. Mathematically, the following set of equations captures this idea for each statement:

$$IN[S] = \text{UNION}_{P \text{ in predecessors of } S} OUT[P]$$

$$OUT[S] = DEF[S] + (IN[S] - KILL[S])$$

This idea can be expressed in Rascal quite literally:

```
public rel[stat, def] reachingDefinitions(
    rel[stat,var] DEFS,
    rel[stat,stat] PRED) {
    set[stat] STATEMENT = carrier(PRED);
    rel[stat,def] DEF = definition(DEFS);
    rel[stat,def] KILL = kill(DEFS);
}
```

```

// The set of mutually recursive dataflow equations
// that has to be solved:

rel[stat,def] IN = {};
rel[stat,def] OUT = DEF;

solve (IN, OUT) {
  IN  = {<S, D> | int S <- STATEMENT,
                stat P <- predecessors(PRED,S),
                def D <- OUT[P]};
  OUT = {<S, D> | int S <- STATEMENT,
                def D <- DEF[S] + (IN[S] - KILL[S])};
};
return IN;
}

```

First, the relations IN and OUT are declared and initialized. Next follows a solve statement that uses IN and OUT as variables and contains two equations that resemble the mathematical equations given above. Note the use of the library function predecessors to obtain the predecessors of a statement for a given control flow graph.

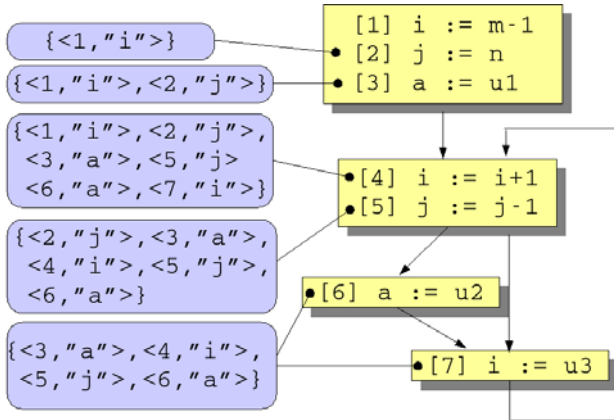


Fig. 1.19. Reaching definitions for example

For our running example the results are as follows (see Figure 1.19, “Reaching definitions for example”). Relation IN has as value:

```

{ <2, <1, "i">>, <3, <2, "j">>, <3, <1, "i">>,
  <4, <3, "a">>, <4, <2, "j">>, <4, <1, "i">>,
  <4, <7, "i">>, <4, <5, "j">>, <4, <6, "a">>,
  <5, <4, "i">>, <5, <3, "a">>, <5, <2, "j">>,
  <5, <5, "j">>, <5, <6, "a">>, <6, <5, "j">>,

```

```

    <6, <4, "i">>, <6, <3, "a">>, <6, <6, "a">>,
    <7, <5, "j">>, <7, <4, "i">>, <7, <3, "a">>,
    <7, <6, "a">>
}

```

If we consider statement 3, then the definitions of variables *i* and *j* from the preceding two statements are still valid. A more interesting case are the definitions that can reach statement 4:

- The definitions of variables *a*, *j* and *i* from, respectively, statements 3, 2 and 1.
- The definition of variable *i* from statement 7 (via the backward control flow path from 7 to 4).
- The definition of variable *j* from statement 5 (via the path 5, 7, 4).
- The definition of variable *a* from statement 6 (via the path 6, 7, 4).

Relation OUT has as value:

```

{ <1, <1, "i">>, <2, <2, "j">>, <2, <1, "i">>,
  <3, <3, "a">>, <3, <2, "j">>, <3, <1, "i">>,
  <4, <4, "i">>, <4, <3, "a">>, <4, <2, "j">>,
  <4, <5, "j">>, <4, <6, "a">>, <5, <5, "j">>,
  <5, <4, "i">>, <5, <3, "a">>, <5, <6, "a">>,
  <6, <6, "a">>, <6, <5, "j">>, <6, <4, "i">>,
  <7, <7, "i">>, <7, <5, "j">>, <7, <3, "a">>,
  <7, <6, "a">>
}

```

Observe, again for statement 4, that all definitions of variable *i* are missing in OUT[4] since they are killed by the definition of *i* in statement 4 itself. Definitions for *a* and *j* are, however, contained in OUT[4]. The result of reaching definitions computation is illustrated in Figure 1.19, “Reaching definitions for example”. We will use the function `reachingDefinitions` later on in Section 5.7, “Program Slicing” when defining program slicing.

5.6.3 Live Variables

The live variables of a statement are those variables whose value will be used by the current statement or some successor of it. The mathematical formulation of this problem is as follows:

$$IN[S] = USE[S] + (OUT[S] - DEF[S])$$

$$OUT[S] = \bigcup_{S' \text{ in successors of } S} IN[S']$$

The first equation says that a variable is live coming into a statement if either it is used before redefinition in that statement or it is live coming out of the statement and is not

redefined in it. The second equation says that a variable is live coming out of a statement if and only if it is live coming into one of its successors.

This can be expressed in Rascal as follows:

```
public rel[stat,def] liveVariables(rel[stat, var] DEFS,
                                   rel[stat, var] USES,
                                   rel[stat,stat] PRED){
    set[stat] STATEMENT = carrier(PRED);
    rel[stat,def] DEF  = definition(DEFS);
    rel[stat,def] USE  = use(USES);

    rel[stat,def] LIN = {};
    rel[stat,def] LOUT = DEF;

    solve(LIN, LOUT) {
        LIN = { <S, D> | stat S <- STATEMENT,
                        def D <- USE[S] +
                                (LOUT[S] - (DEF[S])) };
        LOUT = { <S, D> | stat S <- STATEMENT,
                        stat Succ <- successors(PRED,S),
                        def D <- LIN[Succ] };
    }
    return LIN;
}
```

The results of live variable analysis for our running example are illustrated in Figure 1.20, “Live variables for example”.

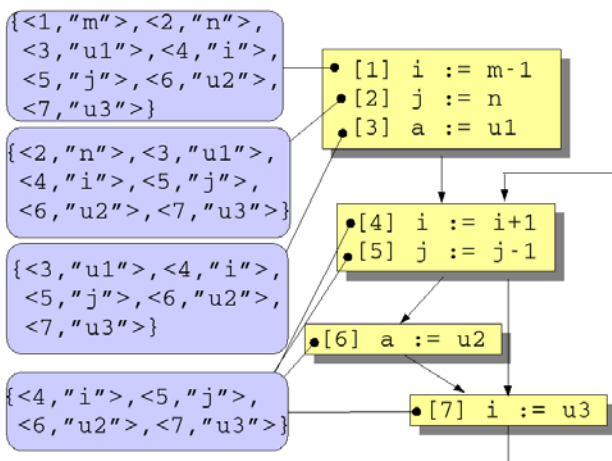


Fig. 1.20. Live variables for example

5.6.4 Wrap Up

The examples in Section 5.6, “Dataflow Analysis” illustrate the following:

- The use of relations to represent control flow and dataflow graphs.
- The use of the solve statement to perform fixed point computations (in these examples to solve dataflow equations).

5.7 Program Slicing

Program slicing is a technique proposed by Weiser [Wei84] for automatically decomposing programs in parts by analyzing their data flow and control flow. Typically, a given statement in a program is selected as the *slicing criterion* and the original program is reduced to an independent subprogram, called a *slice*, that is guaranteed to represent faithfully the behavior of the original program at the slicing criterion. An example will illustrate this (we use line numbers for later reference):

[1] read(n)	[1] read(n)	[1] read(n)
[2] i := 1	[2] i := 1	[2] i := 1
[3] sum := 0	[3] sum := 0	
[4] product := 1		[4] product := 1
[5] while i<= n	[5] while i<= n	[5] while i<= n
do	do	do
begin	begin	begin
[6] sum :=	[6] sum :=	
sum + i	sum + i	
[7] product :=		[7] product :=
product * i		product * i
[8] i := i + 1	[8] i := i + 1	[8] i := i + 1
end	end	end
[9] write(sum)	[9] write(sum)	
[10] write(product)		[10] write(product)
(a) Sample program	(b) Slice for statement [9]	(c) Slice for statement [10]

The initial program is given as (a). The slice with statement [9] as slicing criterion is shown in (b): statements [4] and [7] are irrelevant for computing statement [9] and do not occur in the slice. Similarly, (c) shows the slice with statement [10] as slicing criterion. This particular form of slicing is called *backward slicing*. Slicing can be used for debugging and program understanding, optimization and more. An overview of slicing techniques and applications can be found in [Tip95]. Here we will explore a relational formulation of slicing adapted from a proposal in [JR94]. The basic ingredients of the approach are as follows:

- We assume the relations PRED, DEFS and USES as before.

- We assume an additional set `CONTROL-STATEMENT` that defines which statements are control statements.
- To tie together dataflow and control flow, three auxiliary variables are introduced:
 - The variable `TEST` represents the outcome of a specific test of some conditional statement. The conditional statement defines `TEST` and all statements that are control dependent on this conditional statement will use `TEST`.
 - The variable `EXEC` represents the potential execution dependence of a statement on some conditional statement. The dependent statement defines `EXEC` and an explicit (control) dependence is made between `EXEC` and the corresponding `TEST`.
 - The variable `CONST` represents an arbitrary constant.

The calculation of a (backward) slice now proceeds in six steps:

- Compute the relation `rel[use,def]` *use-def* that relates all uses to their corresponding definitions. The function `reaching-definitions` shown earlier in Section 5.6.2, “Reaching Definitions” does most of the work.
- Compute the relation `rel[def,use]` *def-use-per-stat* that relates the *internal* definitions and uses of a statement.
- Compute the relation `rel[def,use]` *control-dependence* that links all `EXECs` to the corresponding `TESTs`.
- Compute the relation `rel[use,def]` *use-control-def* that combines *use/def* dependencies with control dependencies.
- After these preparations, compute the relation `rel[use,use]` *USE-USE* that contains dependencies of uses on uses.
- The backward slice for a given slicing criterion (a use) is now simply the projection of *USE-USE* for the slicing criterion.

This informal description of backward slicing can now be expressed in Rascal:

```
module demo::Slicing

import Set;
import Relation;
import demo::ReachingDefs;
import demo::Dominators;
import UnitTest;

set[use] BackwardSlice(set[stat] CONTROLSTATEMENT,
                      rel[stat,stat] PRED,
                      rel[stat,var] USES,
```

```

        rel[stat,var] DEFS,
        use Criterion) {

rel[stat, def] REACH = reachingDefinitions(DEFS, PRED);

// Compute the relation between each use and
// corresponding definitions: use_def

rel[use,def] use_def =
{<<S1,V>, <S2,V>> | <stat S1, var V> <- USES,
                   <stat S2, V> <- REACH[S1]};

// Internal dependencies per statement

rel[def,use] def_use_per_stat =
    {<<S,V1>, <S,V2>> | <stat S, var V1> <- DEFS,
                        <S, var V2> <- USES}
    +
    {<<S,V>, <S,"EXEC">> | <stat S, var V> <- DEFS}
    +
    {<<S,"TEST">,<S,V>> | stat S <- CONTROLSTATEMENT,
                        <S, var V> <-
                            domainR(USES, {S})};

// Control dependence: control-dependence

rel[stat, set[stat]] CONTROLDOMINATOR =
domainR(dominators(PRED, 1), CONTROLSTATEMENT);

rel[def,use] control_dependence =
{ <<S2, "EXEC">,<S1,"TEST">>
  | <stat S1, stat S2> <- CONTROLDOMINATOR};

// Control and data dependence: use-control-def

rel[use,def] use_control_def =
    use_def + control_dependence;
rel[use,use] USE_USE =
    (use_control_def o def_use_per_stat)*;

return USE_USE[Criterion];
}

```

Let's apply this to the example from the start of this section and assume the following:

```
rascal> import demo::Slicing;
```

ok

```
rascal> rel[stat,stat] PRED = { <1,2>, <2,3>, <3,4>,
                                <4,5>, <5,6>, <5,9>,
                                <6,7>, <7,8>, <8,5>,
                                <8,9>, <9,10> };

rel[stat,stat]: {<1,2>, ...

rascal> rel[stat,var] DEFS = { <1, "n">, <2, "i">,
                                <3, "sum">,
                                <4,"product">,
                                <6, "sum">,
                                <7, "product">,
                                <8, "i"> };

rel[stat,var]: {<1, "n">, ...

rascal> rel[stat,var] USES = { <5, "i">, <5, "n">,
                                <6, "sum">, <6,"i">,
                                <7, "product">, <7, "i">,
                                <8, "i">, <9, "sum">,
                                <10, "product">
                                };

rel[stat,var]; { <5, "i"> ...

rascal> set[int] CONTROL-STATEMENT = { 5 };
set[int]: {5}

rascal> BackwardSlice(CONTROL-STATEMENT,
                        PRED, USES, DEFS, <9, "sum">);
set[use]: { <1, "EXEC">, <2, "EXEC">, <3, "EXEC">,
            <5, "i">, <5, "n">, <6, "sum">, <6, "i">,
            <6, "EXEC">, <8, "i">, <8, "EXEC">,
            <9, "sum"> }
```

Taking the domain of this result, we get exactly the statements in (b) of the example.

This example illustrates once more the use of relations as the basis for program analysis.

5.8 Visualizing Extracted Information

We are now interested in extracting information from Java source code and in visualizing this information. More precisely:

- Given are Java source files that have been imported to Eclipse as an Eclipse project.
- Extract all class declarations from the sources.

- For each class declaration extract:
 - The number of interfaces implemented by the class.
 - The number of attributes (fields) declared in the class.
 - The number of methods declared in the class.
- Visualize this information as follows:
 - Each class is represented by a rectangle.
 - The color of the rectangle represents the number of implemented interfaces.
 - The width of the rectangle represents the number of fields.
 - The height of the rectangle represents the number of methods.

The Rascal program to achieve this is as follows:

```
module Metrics

import vis::Render;
import vis::Figure;
import Resources;
import JDT;
import Java;
import Set;
import Map;
import IO;

public Figure metrics(loc project) {
    println("Extracting facts from <project>:");
    facts = extractProject(project);
    return visualize(facts);
}

private FProperty popup(str S){
    return mouseOver(box([fillColor("yellow")],
                        text([fontSize(15),
                            fillColor("black")], S)));
}

private Figure visualize(Resource facts) {
    classes =
    {c | c:entity([_*,class(_)]) <- facts@declaredTopTypes};
```

```

fields =
(e : size((facts@declaredFields)[e]) | e <- classes);

methods =
(e : size((facts@declaredMethods)[e]) | e <- classes);

ifaces =
(e : size((facts@implements)[e]) | e <- classes);

sizes = (e : 1.length | <1,e> <- facts@types);

sc = colorScale(toList(ifaces<1>),
                color("grey"), color("red"));
println("<size(classes)> classes\n");

println("Creating visualization ...");
return pack([size(300, 300),gap(10),center()],
            [box([width(2*fields[e]),
                  height(2*methods[e]),
                  fillColor(sc(ifaces[e])),
                  popup("<readable(e)>
                        Fields:<fields[e]>;
                        Methods:<methods[e]>;
                        Interfaces:<ifaces[e]>")
                  ])
              | e <- classes]]);
}

public void main(){
  render(metrics(|project://org.eclipse.imp.pdb.values|));
}

```

It consist of the following parts:

- First relevant libraries are imported, in particular those for visualisation and Java fact extraction.
- Next, the function `metrics` is declared: it takes a location of a Java project, extracts the facts for this project and calls `visualize` to do the actual visualization.
- The function `popup` produces a text popup to be used in the visualization (it will appear when hovering over a figure).
- The heavy lifting is done in the function `visualize`. First, the facts are massaged in a form suitable for display, and then a list of boxes is created with the visual properties corresponding to each class. Finally, these boxes are packed together to achieve a minimal display area for the visualization.

- Function `main`, finally, calls `metrics` for a specific project and renders the resulting figure.

The resulting visualization, embedded in the Rascal IDE, is shown in Figure 1.21, “Rascal IDE showing the Metrics visualization”.




Fig. 1.21. Rascal IDE showing the Metrics visualization

This example illustrates the following:

- Use of the JDT library for fact extraction from Java source code.
- Use of the Figure library for visualization.
- Rascal as bridge between fact extraction and visualization.

6 Concluding Remarks

Rascal and its IDE are in full development at the time of writing and a prototype implementation is available for download. We have given here only a sketch of the language and its applications. The following topics have not been covered in this article:

- The use of SDF modules to parse source text.
- The extensive Rascal library that supports many operations on basic data types including shortest path, reachability and bisimulation on graphs. It also provides tools for drawing graphs and charts, for extracting data from Subversion repositories and more. As mentioned before, our design goal is that Rascal and supporting libraries and tools give as much support as possible for the EASY workflow. This means that we want to provide as much functionality as possible for extraction, analysis, and synthesis. Currently, we provide access to various data and repository formats (XML, HTML, RSF, CVS, SVN, GIT). The Rascal Eclipse JDT library provides direct access to facts that have been extracted from Java source code. The Figure library provides visualization tools. Since the number potentially relevant technologies is unlimited we cannot strive for any form of completeness, but we extend the library on a call by need basis.
- The Rascal IDE that is based on Eclipse and that provides, for instance, very good interactive debugging facilities.

We refer the interested reader to <http://www.rascal-mpl.org/> for a more complete and up-to-date overview and for downloading the latest version of Rascal that extends and enhances Rascal version 0.1 as described here.

The main contribution of this work is providing a language, libraries and tools that make it simpler to carry out software engineering tasks that fit the EASY paradigm. The Rascal language is based on many known concepts but adds some innovations as well: the deep integration of grammars (non-terminals are types; parsers are generated on the fly), integrated pattern matching for regular expressions, abstract and concrete syntactic patterns, rewrite rules as normalization device for structured values, and the visit statement for expressing tree visits illustrate this. We have designed Rascal as the glue that can combine the diverse technologies that are needed to carry out a range of tasks in the domain of software analysis and transformation.

Acknowledgements

Rascal has been designed and implemented by the authors but they have received strong support, encouragement, and help from the following individuals. We are very grateful to them.

Emilie Balland implemented the Rascal debugger in the Eclipse version of Rascal during her visit to CWI in the summer of 2009. Bas Basten provided useful feedback on the design and is developing a Rascal/Eclipse JDT interface that makes it easy to extract facts from Java source code. Joppe Kroon further improved the JDT interface.

Bob Fuhrer's inspiring IMP project motivated us to build Rascal on top of IMP. The PDB was designed and implemented during Jurgen's visit to IBM Research in 2007-2008.

Arnold Lankamp implemented a very efficient version of the Program Data Base (PDB), added a binary streaming format, implemented the SGLR invoker, and currently takes

care of deployment issues. Jeroen van den Bos implemented a Rascal library to access subversion repositories. Mark Hills implemented the `dataTime` data type as well as the static typechecker for Rascal.

Claus Brabrand, Karel Pieterse, Frank Tip and Yaroslav Usenko provided feedback on this paper and suggested several improvements.

We thank the anonymous reviewers for their suggestions for improving this paper.

References

- [ASU86] Aho, A.V., Sethi, R., Ullman, J.D.: *Compilers: Principles, Techniques and Tools*. Addison-Wesley, Reading (1986)
- [AZ05] Anderson, P., Zarins, M.: The CodeSurfer software understanding platform. In: *Proceedings of the 13th International Workshop on Program Comprehension (IWPC 2005)*, pp. 147–148. IEEE, Los Alamitos (2005)
- [BBC+10] Bessey, A., Block, K., Chelf, B., Chou, A., Fulton, B., Hallem, S., Henri-Gros, C., Kamsky, A., McPeak, S., Engler, D.: A few billion lines of code later: using static analysis to find bugs in the real world. *ACM Commun.* 53(2), 66–75 (2010)
- [BKK+98] Borovansky, P., Kirchner, C., Kirchner, H., Moreau, P.-E., Ringeissen, C.: An overview of ELAN. In: Kirchner, C., Kirchner, H. (eds.) *Proceedings of the 2nd International Workshop on Rewriting Logic and its Applications (WRLA 1998)*. *Electronic Notes in Theoretical Computer Science*, vol. 15, pp. 55–70 (1998)
- [BBK+07] Balland, E., Brauner, P., Kopetz, R., Moreau, P.-E., Reilles, A.: Tom: Piggybacking rewriting on Java. In: Baader, F. (ed.) *RTA 2007*. LNCS, vol. 4533, pp. 36–47. Springer, Heidelberg (2007)
- [BPM04] Baxter, I., Pidgeon, P., Mehlich, M.: DMS: Program transformations for practical scalable software evolution. In: *Proceedings of the International Conference on Software Engineering (ICSE 2004)*, pp. 625–634. IEEE, Los Alamitos (2004)
- [BDH+01] van den Brand, M.G.J., van Deursen, A., Heering, J., de Jong, H.A., de Jonge, M., Kuipers, T., Klint, P., Moonen, L., Olivier, P.A., Scheerder, J., Vinju, J.J., Visser, E., Visser, J.: The ASF+SDF Meta-environment: A Component-Based Language Development Environment. In: Wilhelm, R. (ed.) *CC 2001*. LNCS, vol. 2027, pp. 365–370. Springer, Heidelberg (2001)
- [Bey06] Beyer, D.: Relational programming with crocopat. In: *Proceedings of the 28th International Conference on Software Engineering, ICSE 2006*, pp. 807–810. ACM, New York (2006)
- [BNL05] Beyer, D., Noack, A., Lewerentz, C.: Efficient relational calculation for software analysis. *IEEE Trans. Software Engineering* 31(2), 137–149 (2005)
- [BKV03] van den Brand, M.G.J., Klint, P., Vinju, J.J.: Term rewriting with traversal functions. *ACM Transactions on Software Engineering Methodology* 12(2), 152–190 (2003)
- [BKVV08] Bravenboer, M., Kalleberg, K.T., Vermaas, R., Visser, E.: Stratego/XT 0.17. A language and toolset for program transformation. *Science of Computer Programming* 72(1-2), 52–70 (2008)

- [DN04] Dams, D., Namjoshi, K.S.: Orion: High-precision methods for static error analysis of C and C++ programs. Technical report. Bell Labs, Bell Labs Technical Memorandum ITD-04-45263Z (2004)
- [Cor06] Cordy, J.R.: The TXL source transformation language. *Science of Computer Programming* 61(3), 190–210 (2006)
- [FKO98] Feijs, L.M.G., Krikhaar, R., Ommering, R.C.: A relational approach to support software architecture analysis. *Software Practice and Experience* 28(4), 371–400 (1998)
- [Hol08] Holt, R.C.: Grokking software architecture. In: *Proceedings of the 15th Working Conference on Reverse Engineering (WCRE 2008)*, pp. 5–14. IEEE, Los Alamitos (2008)
- [HM03] Hedin, G., Magnusson, E.: The JastAdd system - an aspect-oriented compiler construction system. *Science of Computer Programming*, 37–58 (2003)
- [Joh79] Johnson, S.C.: Lint, a program checker. In: McIlroy, M.D., Kernighan, B.W. (eds.) *Unix Programmer's Manual*, 7th edn., vol. 2B. AT&T Bell Laboratories, Murray Hill (1979)
- [JPJ+90] Jourdan, M., Parigot, D., Julié, C., Durin, O., Le Bellec, C.: Design, implementation and evaluation of the FNC-2 attribute grammar system. In: *Proceedings of the ACM SIGPLAN 1990 Conference on Programming Language Design and Implementation, PLDI 1990*, pp. 209–222. ACM, New York (1990)
- [JR94] Jackson, D.J., Rollins, E.J.: A new model of program dependences for reverse engineering. In: *Proceedings of the 2nd ACM SIGSOFT Symposium on Foundations of Software Engineering. ACM SIGSOFT Software Engineering Notes*, vol. 19, pp. 2–10 (1994)
- [KHR07] Kniesel, G., Hannemann, J., Rho, T.: A comparison of logic-based infrastructures for concern detection and extraction. In: *Proceedings of the 3rd Workshop on Linking Aspect Technology and Evolution, LATE 2007*, p. 6. ACM, New York (2007)
- [Kli93] Klint, P.: A meta-environment for generating programming environments. *ACM Transactions on Software Engineering and Methodology* 2(2), 176–201 (1993)
- [Kli03] Klint, P.: How understanding and restructuring differ from compiling—a rewriting perspective. In: *Proceedings of the 11th International Workshop on Program Comprehension (IWPC 2003)*, pp. 2–12. IEEE Computer Society, Los Alamitos (2003)
- [Kli08] Klint, P.: Using Rscript for software analysis. In: *Working Session on Query Technologies and Applications for Program Comprehension, QTAPC 2008* (2008)
- [KvdSV09] Klint, P., van der Storm, T., Vinju, J.J.: RASCAL: A domain specific language for source code analysis and manipulation. In: *IEEE International Workshop on Source Code Analysis and Manipulation (SCAM 2009)*, pp. 168–177. IEEE Computer Society, Los Alamitos (2009)
- [KvdSV10] Klint, P., van der Storm, T., Vinju, J.: DSL tools: Less maintenance? In: *Preliminary Proceedings of the Tenth Workshop on Language Descriptions Tools and Applications LDTA 2010*, March 27–28 (2010)
- [Kri99] Krikhaar, R.L.: *Software Architecture Reconstruction*. PhD thesis. University of Amsterdam (1999)

- [LR01] Lämmel, R., Riedewald, G.: Prological Language Processing. In: Proceedings of the First Workshop on Language Descriptions, Tools and Applications (LDTA 2001), Genova, Italy, April 7 (2001); van den Brand, M., Parigot, D.: Satellite event of ETAPS 2001. ENTCS, vol. 44. Elsevier Science, Amsterdam (April 2001)
- [McC76] McCabe, T.J.: A complexity measure. *IEEE Transactions on Software Engineering* SE-12(3), 308–320 (1976)
- [dMSV+08] de Moor, O., Sereni, D., Verbaere, M., Hajiyeve, E., Avgustinov, P., Ekman, T., Ongkingco, N., Tibble, J.: QL: Object-oriented queries made easy. In: Lämmel, R., Visser, J., Saraiva, J. (eds.) GTTSE 2007. LNCS, vol. 5235, pp. 78–133. Springer, Heidelberg (2008)
- [MK88] Müller, H., Klashinsky, K.: Rigi – a system for programming-in-the-large. In: Proceedings of the 10th International Conference on Software Engineering (ICSE 10), pp. 80–86 (April 1988)
- [Par07] Parr, T.: The Definitive ANTLR Reference: Building Domain-Specific Languages. Pragmatic Bookshelf (2007)
- [SRK07] Speicher, D., Rho, T., Kniesel, G.: Jtransformer - eine logikbasierte infrastruktur zur codeanalyse. In: Workshop Software-Reengineering, WSR 2007, May 02-04 (2007)
- [Tip95] Tip, F.: A survey of program slicing techniques. *Journal of Programming Languages* 3(3), 121–189 (1995)
- [Wei84] Weiser, M.: Program slicing. *IEEE Transactions on Software Engineering* SE-10(4), 352–357 (1984)
- [VWBGK10] Van Wyk, E., Bodin, D., Gao, J., Krishnan, L.: Silver: An extensible attribute grammar system. *Sci. Comput. Program.* 75(1-2), 39–54 (2010)