UNIVERSITEIT VAN AMSTERDAM

MASTERS PROJECT

Representation Mismatch Reduction for Development in Rules-Based Business Engines

Author:
Paul SPENCER

Supervisor:
Dr. Clemens GRELCK

A thesis submitted in fulfillment of the requirements
for the degree of Master of Software Engineering

in the

Graduate School of Informatics
Faculty of Science

July 25, 2021

UNIVERSITY OF AMSTERDAM

Declaration of Authorship

I, Paul SPENCER, declare that this thesis titled, "Representation Mismatch Reduction for Development in Rules-Based Business Engines" and the work presented in it are my own. I confirm that:

- This work was done wholly or mainly while in candidature for a research degree at this University.

- Where any part of this thesis has previously been submitted for a degree or any other qualification at this University or any other institution, this has been clearly stated.

- Where I have consulted the published work of others, this is always clearly attributed.

- Where I have quoted from the work of others, the source is always given. Except for such quotations, this thesis is entirely my work.

- I have acknowledged all of the main sources of help.

- Where the thesis is based on work done by myself jointly with others, I have made clear exactly what was done by others and what I have contributed myself.

Signed:

Date:

UNIVERSITEIT VAN AMSTERDAM

Abstract

Graduate School of Informatics

Faculty of Science

Master of Software Engineering

Representation Mismatch Reduction for Development in Rules-Based Business Engines

by Paul SPENCER

Context: Declarative rules engine languages, such as Drools, can become difficult to reason about when there are many rules.

Objective: This project investigates how different projections of the code can ease the comprehensibility of the code.

Method: We created an implementation of the Drools language using the MPS language workbench and made innovative projections of large ASTs.

Results:

Keywords: projectional editing; Rules Engines; MPS; Drools

Paper type: Research paper

Acknowledgements

Contents

List of Figures

List of Tables

Chapter 1

Introduction

> The limits of my language mean the
> limits of my world.
>
> ———————————————————————
> *Logico-Tractatus Philosophicus*
> *Ludwig Wittgenstein*

1.1 Problem statement

Miller's Law[1] states that an average human can hold in his short-term memory 5-9 objects. This is often an argument for more succinct code. The argument being anything that is not immediately in the developers vision has to be stored in her memory. With it being impractical to reason about code that she cannot recall, then the fewer relevant items to her reasoning that are out of view the easier it is to reason about the code.

[TODO: complete the problem statement]

1.2 Research questions

To reason about a large code base of rules engine code effectively, a different presentation is needed. This presentation should allow a clearer organization whilst remaining interactive. We can formulate the following research questions based on the discussion in the preceding sections.

The research question we wish to answer is:

- **Main research question:** "How can projectional editors and DSLs be combined to address feedback mechanisms for developers in the context of reasoning about rules in a rule-based business engine?"

This question requires knowing if it is possible with current tooling, thus we would like to answer the question:

- **RQ 1:** "What is the current state of language workbenches supporting projectional editing?"

Finally, we specifically would like to know how we can improve the ability to reason about the business rules engine, so we ask the question:

- **RQ 2:** "Which projections can help developers to get appropriate feedback about rules?"

1.3   Contributions

This thesis proposes a code representation of business rules in a concise and readable format that could solve comprehensibility issues resulting from large code bases of business rules. The implementation behind the approach relies on language engineering and projectional editing. An implementation has been developed as a stand alone opensource solution on a limited demonstration version of Drools. The underlying Drools implementation can be used as a base language for model to model generation by the wider MPS ecosystem.

1.4   Project context

This investigation was hosted by Khonraad Software Engineering, a subsidiary of Visma. Khonraad provides mission-critical services focussed on the automation of workflows at the cross-section of local government and healthcare. Specifically, Khonraad facilitates the mental health care and coercion laws in the Netherlands - WVGGZ, WZD, and WTH - which provide agencies the ability to intervene in domestic violence, psychiatric disorders, and illnesses.

Khonraad's system facilitates reporting and communication between municipalities, police, judiciary, lawyers, mental health care, and many social care institutions. The system has 15,000 users and is available 24/7.

Configuration and administration use complex matrices of compliance mechanisms, access user rights and communication settings. The sensitivity of the personal data, being both medical and criminal, means security is of utmost importance. The security against data loss, preventing unlawful disclosure and guaranteeing availability, especially during crisis situations, is crucial. Demonstration of the correctness of the, often changing, configuration is a major concern in the company.

This work environment allows us to work on an existing project, where the tangible success will have an impact on the lives of those in critical need. Khonraad has it's own implementations in the Drools language, that have evolved over the iterations of the laws. The evolution of the code base over the years means that the real-life issues we came across are not just thought experiments.

1.5   Thesis outline

We start in chapter 2 with the required background information on projectional editing and rules engines. In chapter 3 we present the research questions. Further, the chapter

describes the protocol that we use for search strategy, selecting our studies, extracting data from them, and synthesizing the results. Chapter 4 presents the results of our synthesis of data from the primary studies. This is followed, in chapter 5, by a discussion of both the validity of the work and the implications of the findings. We discuss the implications of this study in chapter 6. Finally, the conclusions are presented in chapter 7.

Chapter 2

Background

This chapter gives the background information required on rules engines and projectional editing. It presents the specific case of rules engine that we will be using for our investigation: Drools. Further, it briefly examines the base tool type for creating Domain-specific languages: Language work benches. Finally, it presents the specific projectional editing tool we will be using: JetBrains MPS.

2.1 RulesEngines

2.1.1 What is a rules engine?

In this section we will describe what a rules engine is and a little of its history.

The Aristotelian doctrine of essentialism declares that a thing has properties that are essential and properties that are accidental. If one takes away accidental properties, then the thing remains the thing. If one takes away essential properties, the thing is no longer the thing. If the thing is a business application, then its essential properties are its business rules.

Simply put, business rules are the principles or regulations by which an organization carries out the tasks needed to achieve their goals. When properly defined these rules can be encoded into statements that defines or constrains some aspect of the business organizational behaviour. A rule consists of a condition and an action. When the condition is satisfied then the action is performed. More formally, business rules can be seen as the implication in the basic logical principle of Modus Ponens.

When described like this, one could me forgiven for thinking is this not just an if-then logic that is frequently used in traditional programming. One would not be wrong, however in traditional programming, representing all the combinatorial outcomes can become complex. In the typical application architecture, rules are distributed in the source code or database. Each additional rule leads to more fragility.

Documentation describing these rules may be found in the design documentation or user manuals. However, as applications evolve documentation gets out of sync with codebase. Once this desynchronization occurs, to know what the rules that govern the application, one has to navigate the codebase and decode the rules from their, often scattered, locations.

A rules engine is also known as a Business Rules Engine, a Business Rules Management System or a Production Rules System. The goal of a rules engine is the abstraction of business rules into encoded and packaged logic that defines the tasks of an organization with the accompanying tools that evaluate and execute these rules. Simply put, they are where we evaluate our rules. Rules engines match rules against facts and infer conclusions. If we return to the Modus Ponens comparison:

$$\frac{\begin{array}{c} p \\ p \to q \end{array}}{\therefore q}$$

If the premise p holds. And the implication $p \to q$ holds then the conclusion q holds. In terms of a rule engine and business rules this could be seen as:

1. the rules engine gathers the data for the premise: p

2. it examines the business rules as the implications: $p \to q$

3. it executes the conclusion: q

Rules Engines are declarative, focussing on the what of the rules not the the how of the execution. Date[2] describes rules engine as to "specify business process declaratively, via business rules and get the system to compile those rules in to the necessary procedural (and executable) code." Fowler[3] describes rules engine as follows: " ... providing an alternative computational model. Instead of the usual imperative model, which consists of commands in sequence with conditionals and loops, a rules engine is based on a Production Rule System. This is a set of production rules, each of which has a condition and an action ...".

Rule engines arose from the expert systems of the late 70s and early 80s. Expert systems initially had three main techniques for knowledge representation: Rules, frames and logic[4]. "The granddaddy" of the expert systems, MYCIN, relied heavily on rules based knowledge representation[5], rather than long inference chains. MYCIN was used to identify bacteria and recommend antibiotic prescriptions. MYCIN and its progenitor, DENDRAL, spawned a whole family of Clinical Decision Support Systems that pushed the rules engine technology until the early 1980's. Research into rules engines died out in the 1980s as it fell out of fashion.

Early in their existence, the rules engines hit a limiting factor because the matching algorithms they used suffered from the utility problem, i.e. the match cost increased linearly with the number of rules being examined/ This problem was solved by Charles Forgy's efficient pattern matching Rete algorithm[6], and its successors. This algorithm works by modelling the rules as a network of nodes where each node type works as a filter. A fact will be filtered through this network. The pre-calculation of this network is what provides the performance characteristics.

The first popular rules engine was Office Production System from 1976. In 1981 OPS5 added the Rete algorithm. CLIPS in .. JESS Drools

In general, rules engines are forward chaining. This means to test if [TODO: Explain forward chaining with logic symbols]

[TODO: ADD MORE HISTORY HERE]

Product		Developer	licence type
CLIPS	[7]	NASA	open source
Drools	[8]	JBoss/RedHat	open source
BizTalk Business Rule Engine	[9]	Microsoft	proprietary
WebSphere ILOG JRules	[10]	IBM	proprietary
OpenRules	[11]	OpenRules	open source

TABLE 2.1: Rules Engine products.

Moving forward to current times, there are a few rules engines currently in use. Some of the more commonly used ones are shown in table 2.1

Some of the advantages of using a rules engine include:

- The separation of knowledge from it's implementation logic

- Business logic can be externalized

- Rules can be human readable

Rules that represent policies are easily communicated and understood. Rules retain a higher level of independence than conventional programming languages. Rules separate knowledge from its implementation logic. Rules can be changed without changing source code; thus, there is no need to recompile the application's code. Cost of production and maintenance decreases.

In summary a rules engine, is the executor of a rules based program, consisting of discreet declarative rules which model a part of the business domain.

2.1.2   What is Drools?

JBoss Rules, or as it is more commonly known, Drools, is the leading opensource rules engine written in Java. In this paper when we use the name "Drool" we are referring to the "Drools Expert" which is the rule engine module of the Drools Suite. Drools started in 2001, but rose to prominence with it's 2005 2.0 release. It is an advanced inference engine using an enhanced version of the Rete algorithm, called ReteOO[12], adapted to an object-oriented interface specifically for Java. Designed to accept pluggable language implementations, it can also work with Python and .Net. It is considered one of the most developed and supported rules platforms.

For rules to be executed there are 4 major components as demonstrated in figure 2.1. The production memory contains the rules. This will not change during an analysis session. The rules are the focus of this thesis and therefore we will delve into much more detail later on these.

In Forgy's[6] overview of a rete algorithm, the following steps occur.

1. Match : Evaluate the LHSs of the productions to determine which are satisfied given the current contents of working memory

2. Conflict resolution : Select one production with a satisfied LHS; if no productions have satisfied LHSs, halt the interpreter

3. Act : Perform the actions in the RHS of the selected production

4. Re-evaluate : Go To 1

Figure 2.2 show more detail of how these components interact within Drools to infer a conclusion. First a fact or facts are asserted in the working memory. The working memory contains the current state of the facts. This triggers the inference engine. The pattern matcher, using the aforementioned ReteOO algorithm will determine examine the working memory and a representation of the rules from the production memory to determine which rules are true. Matching rules will be placed on the agenda. It can be the case that many rules are concurrently true for the same fact assertion. These rules are in conflict. A conflict resolution strategy will decide which rule will fire in which order from the agenda. The first rule on the agenda will fire. If the rule modifies, retracts or asserts a fact, then the inference loop begins again. If a rule specifies to halt or there are no matching rules left on the agenda, we have inferred our conclusion.
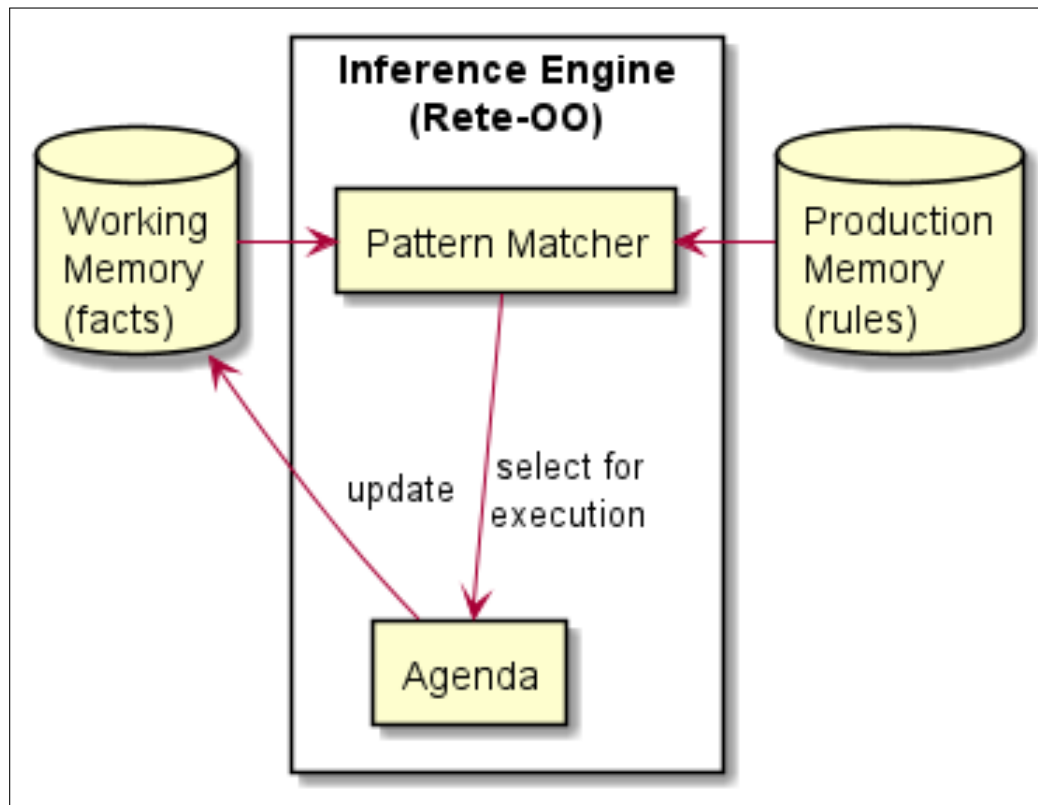
FIGURE 2.1: Drools components.

The component we will be focussing on in this paper is the rules. Rules are stored in a rules file, a text file, typically with a drl extension. During execution the rules do not change and are stored in production memory. For the sake of this paper we will skip past package, import, global, declare, function and query, which are also stored in the rule file. We will examine the anatomy of a rule.

A rule is made of 3 parts: attributes; conditions; and consequences. Attributes are an optional hints to the inference engine as to how the rule should be examined. The

FIGURE 2.2: Drools Inference Loop.

conditional, when, or left hand side (LHS) of the rule statement is a block of conditions that have to in aggregate return true for the asserted fact in order to be considered to be placed on the agenda. The actions, consequences, then, or Right hand side (RHS) of the rule statement contains actions to be executed, should the rule be filtered

The LHS is a predicate statement, made up of a number of patterns. variables can be bound to facts that match these patterns for use later in the LHS or for updating the working memory on the RHS. 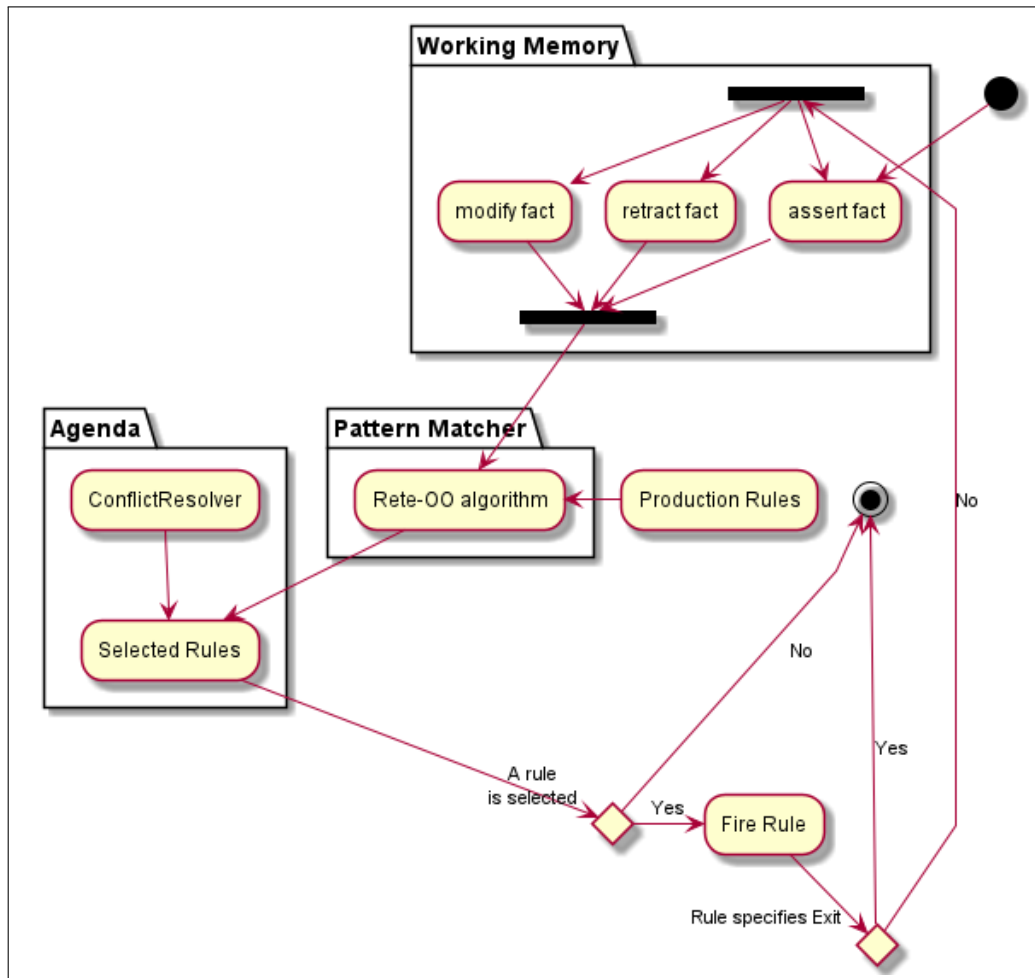the patterns are used to evaluate against the working memory. The pattern match against the existence of facts. Patterns can also match against conditions of the properties of facts. Connectives such as not, and, and or can be applied to the patterns. The patterns apply to individual Facts rather than the group, thus can be seen as first order predicates.

There are some more advanced features in the LHS, but for this paper, these are the features we will be looking at.

Whilst the RHS can contain arbitrary code to be executed when a rule is fired, it's main purpose is to adjust the state of truth in the working memory. One can insert, modify, and retract facts in the working memory. modifying and retracting facts, must be done on fact variable references that have been created in the LHS. One can explicitly terminate the inference loop, with a halt command.

FIGURE 2.3: Drools Rule Breakdown.

An explanatory example

An example of a DRL file can be seen in listing 2.1. This has been extracted from the Drools sample code.

```
1    package org.drools.examples.honestpolitician
2
3    import org.drools.examples.honestpolitician.Politician;
4    import org.drools.examples.honest politician.Hope;
5
6    rule "We have an honest Politician"
7        salience 10
8        when
9            exists( Politician( honest == true ) )
10       then
11           insertLogical( new Hope() );
12   end
13
14   rule "Hope Lives"
15       salience 10
16       when
17           exists( Hope() )
18       then
19           System.out.println("Hurrah!!! Democracy Lives");
20   end
21
22   rule "Hope is Dead"
23       when
24           not( Hope() )
25       then
26           System.out.println( "We are all Doomed!!! Democracy is Dead
                 " );
27   end
28
29   rule "Corrupt the Honest"
30       when
31           $p : Politician( honest == true )
32           exists( Hope() )
```

```
33          then
34              System.out.println( "I'm an evil corporation and I have
                    corrupted " + $p.getName() );
35              modify( $p ) {
36                  setHonest( false )
37              }
38      end
```

LISTING 2.1: Example Drools file.

Listing 2.1 gives the Drools engine instructions on what actions to take when something changes in the working memory. What this toy example does is reacts to when an honest politician is added to the working memory, prints a message celebrating the existence of said politician, corrupts her, gloats in a message and then prints a message of despair. The code in listing 2.1 does the following:

1. on line 1 the package statement identifies the rule file

2. on lines 3 and 4 the import statements describes which facts can be used

3. the "We have an honest Politician" rule on line 6 does the following:

   (a) using salience on line 7 it sets that this rule is to be run before rules with a lower salience

   (b) on line 10 it checks the working memory for Politician facts with the honest property equal to true

   (c) on line 12, if found then Hope facts will be inserted into the working memory

4. the "Hope Lives" rule on line 15 does the following:

   (a) line 18 check if any Hope facts exist

   (b) on line 20, if found, it prints a message

5. the "Hope is Dead" rule on line 23 does the following:

   (a) checks if no Hope facts exist on line 25

   (b) if none are found, on line 27, it prints a message

6. the "Corrupt the Honest" rule on line 30 does the following:

   (a) line 32 checks for any Politician facts with the honest property equal to true, and sets them to the variable $p

   (b) line 33 checks if any Hope facts exist

   (c) if both hope and politicians are found on line 35 it prints a message including the $p variables name

   (d) on line 36 to 38 it modifies the fact in working memory represented by $p to change it's honest property

2.2 Projectional Editing

2.2.1 What is projectional editing?

2.2.2 what are Language Workbenches?

2.2.3 What is MPS?

Chapter 3

Method

Chapter 4

Results

the purpose of abstraction is not to be vague but to create a new semantic level in which one can be absolutly precise.

Logico-Tractatus Philosophicus
Edsger W. Dijkstra

Chapter 5

Discussion

5.1 Threats to Validity

5.1.1 Construct Validity

5.1.2 Internal Validity

5.1.3 External Validity

5.1.4 Reliability

5.1.5 Repeatability vs Reproducibility

5.1.6 Method improvement

Chapter 6

Implications to research and practice

6.1   Implications to research

6.2   Future research directions

6.3   Implications to practice

Chapter 7

Conclusion

Appendix A

InterviewTranscripts

Write your appendix content here.

Bibliography

[1]   G. A. Miller, "The magical number seven, plus or minus two: Some limits on our capacity for processing information.," *Psychological review*, vol. 63, no. 2, p. 81, 1956.

[2]   C. J. Date, *What not how: the business rules approach to application development*. Addison-Wesley Professional, 2000.

[3]   M. Fowler, *Should i use a rules engine?* `https://martinfowler.com/bliki/RulesEngine.html`, Accessed: 2021-07-18, 2009.

[4]   P. Jackson, "Introduction to expert systems," 1986.

[5]   E. H. Shortliffe, "Mycin: A rule-based computer program for advising physicians regarding antimicrobial therapy selection.," STANFORD UNIV CALIF DEPT OF COMPUTER SCIENCE, Tech. Rep., 1974.

[6]   C. L. Forgy, "Rete: A fast algorithm for the many pattern/many object pattern match problem," in *Readings in Artificial Intelligence and Databases*, Elsevier, 1989, pp. 547–559.

[7]   *CLIPS product page*, `http://www.clipsrules.net/`, Accessed: 2021-07-17.

[8]   *Drools product page*, `https://www.drools.org/`, Accessed: 2021-07-17.

[9]   *BizTalk product page*, `https://docs.microsoft.com/en-gb/biztalk/`, Accessed: 2021-07-17.

[10]   *IBM WebSphere JRules product page*, `https://www.ibm.com/docs/en/iis/11.7?topic=applications-websphere-ilog-jrules`, Accessed: 2021-07-17.

[11]   *OpenRules product page*, `https://openrules.com/`, Accessed: 2021-07-17.

[12]   D. Sottara, P. Mello, and M. Proctor, "A configurable rete-oo engine for reasoning with different types of imperfect information," *IEEE Transactions on Knowledge and Data Engineering*, vol. 22, no. 11, pp. 1535–1548, 2010.